

# Blockchains: A Systematic Multivocal Literature Review

BERT-JAN BUTIJN, Tilburg University, The Netherlands

DAMIAN A. TAMBURRI, Eindhoven University of Technology, The Netherlands

WILLEM-JAN VAN DEN HEUVEL, Tilburg University, The Netherlands

---

Blockchain technology has gained tremendous popularity both in practice and academia. The goal of this article is to develop a coherent overview of the state of the art in blockchain technology, using a *systematic* (i.e., protocol-based, replicable), *multivocal* (i.e., featuring both white and grey literature alike) literature review to (1) define blockchain technology, (2) elaborate on its architecture options and (3) tradeoffs, as well as to understand (4) the current applications and challenges, as evident from the state of the art. We derive a systematic definition of blockchain technology, based on a formal concept analysis. Further, we flesh out an overview of blockchain technology elaborated by means of Grounded-Theory.

CCS Concepts: • General and reference → Surveys and overviews; Empirical studies; • Human-centered computing → Collaborative and social computing theory, concepts and paradigms; • Computing methodologies → Parallel computing methodologies; • Software and its engineering → Designing software;

Additional Key Words and Phrases: Blockchain technology, distributed ledger technology, smart contracts, software architecture, multi vocal literature review, literature review

**ACM Reference format:**

Bert-Jan Butijn, Damian A. Tamburri, and Willem-Jan van den Heuvel. 2020. Blockchains: A Systematic Multivocal Literature Review. *ACM Comput. Surv.* 53, 3, Article 61 (June 2020), 37 pages.

<https://doi.org/10.1145/3369052>

---

## 1 INTRODUCTION

In 2008, Satoshi Nakamoto, a person or group of people,<sup>1</sup> introduced the concept of a peer-to-peer (P2P) version of electronic cash that allows for online payments to be made directly from one party to another without any trusted financial institution [66]. The Nakamoto article preludes the rise of Bitcoin, ushering in the dawn of blockchain technology (BCT). BCT utilizes the concept of a digital distributed ledger (i.e., a record or “book” of transactions), which enables the participants of a P2P network to record transactions that are publicly verifiable. In essence, BCT ensures trust between parties without any trusted intermediary when performing transactions [85]. BCT

---

<sup>1</sup>To this day the identity or identities of Satoshi Nakamoto remains unknown [90].

Some of the authors’ work is supported by the EU H2020 framework programme, grant “ANITA” under Grant No. 787061 and grant “PRoTECT” under Grant No. 815356.

Authors’ addresses: B.-J. Butijn, W.-J. van den Heuvel, Tilburg University, Warandelaan 2, Tilburg, 5037 AB, The Netherlands; emails: b.j.butijn@uvt.nl, W.J.A.M.v.d.Heuvel@jads.nl; D. A. Tamburri, Eindhoven University of Technology, De Rondom 70, 5612 AP, Eindhoven, The Netherlands; email: d.a.tamburri@uvt.nl.

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than ACM must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from [permissions@acm.org](mailto:permissions@acm.org).

© 2020 Association for Computing Machinery.

0360-0300/2020/06-ART61 \$15.00

<https://doi.org/10.1145/3369052>

has gained considerable scholarly attention [91]. For this very reason, we seize the opportunity to conduct and report a systematic and *multivocal* study of the state of the art in BCT for the benefit of further well-founded research as well as practice. As said, we operate a multivocal systematic study, namely, we not only focus on research literature but also consider the so-called *grey* literature (i.e., books, technical reports, whitepapers, and more), which may carry important information concerning the BCT software architecture landscape [72].

We flesh out the results of our study starting from a rigorous definition of what BCT is and is not. Indeed, the terms BCT and distributed ledger technology (DLT) are frequently used interchangeably despite attempts to semantically discern them on their distinctive underlying architectures [35]. Second, using the well-known 4+1 software architecture framework by Kruchten [47] as a lens for analysis, the study outlines the design options available in the state of the art for BCT architectures.

Third, through the scenario perspective several applications are presented to illustrate how BCT could be harnessed for different scenarios. By categorizing and summarizing the current applications of BCT practitioners gain more insight in the rich palette of possibilities BCT has to offer. Fourth, with this study we highlight the properties of BCT and elucidate their arising tradeoffs. Insights on these properties and their tradeoffs aids practitioners in making design choices while providing scholars means to assess blockchain architecture.

Fifth, the research delineates a comprehensive and data-driven overview of the challenges in the field of BCT. Following, this study thoroughly discusses the relation amongst BCT concepts based on a rigorous and systematic analysis of the literature. Establishing this relations can help practitioners further develop BCT while scholars are provided with more accurate measurements to gauge its benefits. The discussion is further strengthened by presenting highlights and observations in-depth attained from the papers under review. Finally, this research presents a systematic overview of all current research gaps to help direct future research endeavors.

## 1.1 Related Work

Several previous related surveys exist in the state of the art, even though none of them have the scope, breadth, and width we adopt in our research design. We report the most closely related here below and highlight the novelty of our work. Yli-Huumo et al. [91] review and map the extant literature to indicate research gaps. This review, however, predominantly features literature related to Bitcoin and corresponding issues. On one hand, the synthesis operated by Yli-Huumo et al. is only loosely systematic, and, on the other hand, the field of BCT has been rapidly developing since the publication of their study [91]. By comparison, our work also fleshes out the architecture of other blockchain networks, and provides a more up-to-date and systematic overview of BCT developments.

In Reference [80] a literature review on cryptocurrencies is presented, which is, however, not focused on BCT architectures. Moreover, in this work we focus not only on the cryptomarket but consider use-cases for blockchain other than cryptocurrency. The same issue recurs with other reviews focused on the literature related to smart contracts, i.e., programs that can be deployed and run on a blockchain [2] and their applications [57]. For example, Bartoletti and Pompianu [7] focus on smart contract applications, and review the platforms and design patterns for such smart contracts. In our own work, we build from these foundations and include a comparison of smart contracts with the architecture principles of BCT.

Furthermore, there are several literature reviews that examine the use-cases and specific applications of BCT, e.g., Karafloski et al. [39]. Several other sector- and domain-specific reviews also exist, e.g., for the Internet of Things (IoT) [18, 41], or Multi-Agent Systems (MAS) [14]. While other papers review how BCT could be utilized for the aforementioned domains, these reviews do

not present BCT applications based on a rigorous scientific evaluation—in our work we set out to operate a systematic synthesis of architecture elements as well as the alternatives in architecture decision-making spanning multiple domains and encompassing reference literature from, e.g., Supply-chain management [44, 84], usage of BCT by governments [9], BCT in healthcare [48], and more.

By contrast, the work of Tama et al. [78] presents a brief critical review of BCT and some applications for multiple fields. Our work, however, provides a more elaborate overview of BCT applications for these fields, and in addition an in-depth insight into the architecture of blockchain technology that has been obtained through a Grounded-Theory approach. Furthermore, our study includes a definition of BCT based on formal concept analysis. What is more, Casino et al. [15] present a review of BCT as a basis for multi-purpose applications design. Rather than concentrating solely on BCT applications, our work in addition provides a rigorous definition of blockchain based on formal concept analysis and offers an extensive multi-vocal catalogue and accompanying descriptions of anything that was published about blockchain technology. The catalogue presented in this article follows the well-known 4+1 view framework [47] for architecture description to aid anyone in framing, operating, deciding upon or describing blockchain architectures in general. Hence, the scope of our work is much broader, not only discussing applications of blockchain but in addition its architecture in depth. Finally, we provide a data-driven, in-depth, and evidence-based overview of research gaps in the field of BCT.

## 1.2 Structure of the Article

The remainder of this article is organized as follows: The next section reviews the background and basic notions of BCT. The methodology section (Section 3) elaborates on the approach taken to attain the results of this study. In the section thereafter (Section 4) a definition of BCT that has been constructed based on the literature is presented. Following, the study presents BCT software architecture from multiple perspectives in Section 5. In Section 6 reviews the scenarios for using BCT found in the literature to provide insight into the applications for which BCT is used. The characteristics of BCT and architectural tradeoffs are presented in Section 7. Challenges of BCT are presented in Section 8, along with an outlook. The results of the research are discussed thereafter (Section 9). Based on the discussion of the results, in Section 10 suggestions for future research are presented. In Section 11 the limitations of this research and potential threats to validity are addressed. Section 12 concludes the article.

## 2 BACKGROUND AND BASIC NOTIONS

As previously stated, BCT first appeared in 2008, featured in the seminal paper “Bitcoin: A Peer-to-Peer Electronic Cash System” by Nakamoto [66]. The paper proposes a P2P electronic cash system that allowed for the execution of transactions between one party and another without requiring a trusted third party to act as a safeguard and check the validity of the transaction. A year later in 2009, the Bitcoin network was launched [94].

### 2.1 Historical Setting

The first solution that Nakamoto suggested to enable the transactions of digital coins is that owners of a coin wishing to commit a transaction should digitally sign a hash of the previous transaction and the public key of the next owner, both is added to the end of the coin. An electronic coin as such is defined as a chain of digital signatures. By verifying the signatures of a coin the payee can verify the historical chain of ownership. However, this provides a payee with no guarantee that the coin has not already been double spent as there is no way to verify that the previous owners did not sign any earlier transactions. *Double-spending* refers to spending the same currency in

two distinct transactions at the same time. In traditional settings, a centralized trusted third party (e.g., bank or mint) verifies whether the owner of a coin did not double spend the same coin. To verify transactions traditional trusted third parties maintain a centralized ledger that records all transactions and the order in which they were enacted. Moreover, the trusted third party needs to be aware of all transactions as there is no other way to confirm the absence of a transaction.

## 2.2 Terms and Definitions

In order for transactions to be executed without a trusted third party there also needs to be full awareness, and a single history of these transactions. In the Bitcoin paper, two solutions are proposed to accomplish the aforementioned goals: (1) Transactions should be publicly announced to all participants in the network. These objectives are attained by employing a distributed ledger on a P2P network. Specific network participants called *nodes* each store a local copy of the ledger. (2) Nodes need to reach a consensus about the history of the transactions, and the order in which they were received. This raises another problem, however: Some of the nodes in the network might behave maliciously and try to change the communication contents. In the literature, this problem is referred to as the *Byzantine Generals Problem* [49]. Non-malicious nodes need to be able to distinguish the information that has been tampered with from the correct information by reaching a consensus over the consistency of the distributed ledger to determine the validity of a transaction. Consequently, this requires proof that when the transaction was executed, the majority of nodes have reached a trustworthy consensus that it was the first received. In essence, these requirements are introduced to ensure that the system to a certain extent can tolerate malicious behavior by nodes participating in the network, to which is commonly referred to as *Byzantine fault tolerance* [71].

In the seminal Bitcoin paper [66] several concepts are presented to satisfy these requirements starting with the use of a timestamp server. The server takes a hash of the block of transactions that are required to be timestamped and publishes the hash on the network. What the *timestamp* proofs is that the data existed at a certain point in time. The hash is *chained* to the previous hash, because the latter time stamp is included in the former. As a result, each additional hash reinforces the ones before it, and as more blocks are added the chain will grow ever stronger. The next concept presented enables nodes to reach a consensus on whether the distributed ledgers are consistent with one another, thus that all transactions are valid. A naive way of accomplishing this would be to let the majority of nodes vote over its consistency. However, that would make the blockchain prone to *Sybil attacks* whereby a malicious attacker creates or copies multiple identities to control the network.

Bitcoin diminishes the possibility of a Sybil attack by employing a *Proof-of-Work (PoW) consensus protocol*, which stipulates that not the majority of IP-addresses count as the majority vote of the network but rather the majority of computational power. While it might be easy for an attacker to create several nodes in a network, amassing large amounts of computational power might prove to be more difficult. The PoW consensus algorithm distributes accounting rights and rewards through a computing power competition in which all nodes of the network can participate. Nodes try to be the first to solve a computational hard mathematical puzzle by finding the right *nonce* (a random number) for the block-header based on information of the prior block. This process is called *mining* and the nodes executing the calculations are referred to as *miners* in the Bitcoin nomenclature. The first miner to finish creates the next block and is rewarded by receiving an amount of Bitcoin. However, because the mining process is probabilistic two or more blocks might be created and propagated by distinct miners simultaneously. These phenomena are known as *forks*. In the event of a fork, nodes as a rule always trust the *longest chain* of blocks as the chain holding the truth with regard to transaction validity (which is analogous to the most computational work). Other nodes

wait until new blocks are proposed after the occurrence of the fork to determine which chain will become the longest chain. Consequently, transactions are not confirmed before a longest chain has formed.

The *longest chain rule* is a safeguard to secure the blockchain against the possibility to delay the propagation of transactions, which in turn, opens the possibility of introducing fake transactions. As the computational power and interests of the miners might vary the PoW consensus protocol increases or decreases the difficulty of the mathematical problem in such a way that the interval between the generation of new blocks, referred to as *block interval time* remains constant at 10 minutes. Tampering with the transactions recorded on the Bitcoin blockchain would therefore require an attacker not only to be the first one to generate the latest block, but also to control the longest chain.

After the introduction of the genesis PoW-based consensus protocol for Bitcoin many others have been introduced for blockchain such as the following: (1) *Proof-of-Stake* (PoS), which replaces PoW-based mining with a mechanism that makes the chances of mining a block proportional to the amount of stake (currency) a miner has [29, 52, 77, 92]; (2) *Delegated-Proof-of-Stake* (DPoS), where the chances of mining a block are also based on a miner stake but allows for the delegation of voting on the correctness of a block [64, 93]; (3) *Proof-of-Elapsed-Time* (PoET), which uses dedicated hardware to create consensus [61, 90]; and (4) *Zero Knowledge Proofs* (ZKP), which aims to provide users performing transactions with more privacy [35, 89, 92]. The advent of blockchain also revived the interest in preexisting consensus protocols such as *Practical Byzantine Fault Tolerance* (PBFT), which could be utilized for a similar purpose as ZKPs [25, 89, 93].

Furthermore, Bitcoin was envisioned as a *public blockchain network* that anyone willing can access, and that is *permissionless*, meaning that everyone connected to the network can request transactions or become a miner to check the validity of transactions. By contrast, in the past decade *private blockchain networks* have been introduced that allow only selected participants from one organization to join the network, which can also perform only actions that are *permissioned* on the network. Finally, *Consortium blockchain networks* can be considered a hybrid approach as the number of participants that can join the network is restricted, but they can be from different organizations. Among the connected participants, the permissions they are granted on the network might differ [74, 89, 93].

Although the Bitcoin technology introduced the concept of BCT to allow for electronic payments using *cryptocurrency* (digital coins) between anonymous peers, nowadays other blockchain networks such as Ethereum offer to possibility to deploy *smart contracts*, that is, programs that can be deployed, run, and verified correct over a blockchain. Smart contracts use triggers, conditions and business logic to enable more complex programmable transactions [89] for the automation of (business) processes [27, 29, 75].

### 3 RESEARCH METHODOLOGY

To attain our results, we conducted a systematic Multivocal Literature Review (MLR) on blockchain technology. Specifically, we address the following research questions:

- RQ1 How can blockchain technology be systematically defined?
- RQ2 What applications of blockchain technology have currently been published and how can these applications be classified?
- RQ3 What are the properties of blockchain technology and what are their tradeoffs?
- RQ4 What are the challenges for blockchain technology?
- RQ5 What are the current research gaps in the field of blockchain technology?

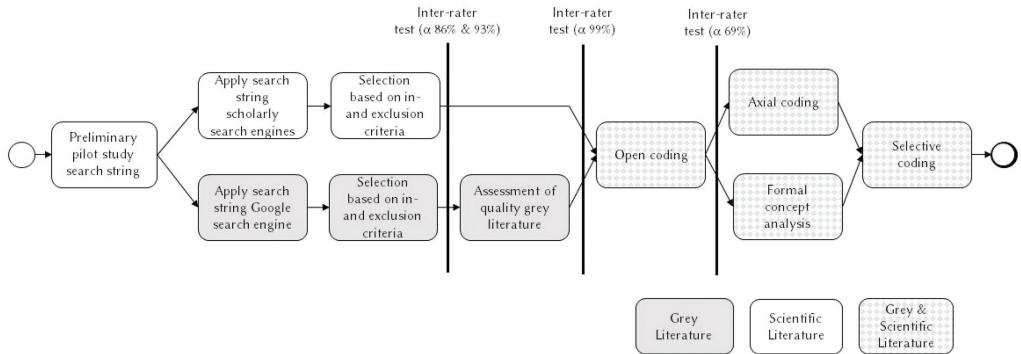


Fig. 1. Research methodology: An outline.

The first research question rotates around providing a systematically-derived definition for blockchains while RQ2 focuses on the applications for which BCT is utilized. RQ3 seeks to offer an overview of the notable properties of BCT (architectural or otherwise) as well as their tradeoffs. The fourth research question focuses on delineating the challenges in the field of BCT. Finally, RQ5 aims at presenting research gaps that future research endeavors can address.

### 3.1 Data Preparation Approach

The benefit of a MLR approach is that, beyond typical systematic literature reviews [43] (SLRs), which use academic peer-reviewed articles alone, a multivocal literature review (MLR) also allows for the inclusion of Grey Literature (GL). GL is typically produced by practitioners, such as private industry, governments, academics and industry, and any party that is not controlled by commercial publishers or peer-review. Generally, therefore, grey literature is not published in books or scientific journals. However, this literature can provide invaluable insights into the state of the practice in a field [28]. Given that at the moment of writing the field of BCT is still relatively in its infancy, we therefore deem it appropriate to include relevant literature created by practitioners in the field of BCT for a better understanding of the field. Including GL in our review allows us to combine and synthesize academic literature with the state-of-the-art in practice.

In conducting our MLR we set out to identify (a) all relevant academic peer-reviewed articles (scientific literature) and (b) all relevant grey literature for this study. To reduce the possibility of researcher bias, a predefined protocol for the identification of both the relevant scientific literature (SL) and grey literature (GL) needs to be established [28]. While carrying out our systematic literature review we followed three steps: (1) create a selection of articles to review, (2) conduct the review, and (3) analyze the data. A process model of the methodology used for this research is depicted in Figure 1.

### 3.2 Search Strategy

The first step has been carried out using the protocol for systematic reviews suggested by Kitchenham [43]. The protocol suggests three stages for a literature review: (1) elaborate the search string, (2) apply the string on chosen search engines; and (3) filter out and extract primary papers based on pre-established exclusion criteria from search results. The implementation of these steps is presented in Figure 2.

The search string was determined by deriving relevant keywords from the research questions. Before carrying out our systematic search, we conducted a preliminary pilot study by experimenting with the search terms to select more results. This process yielded the following search terms:

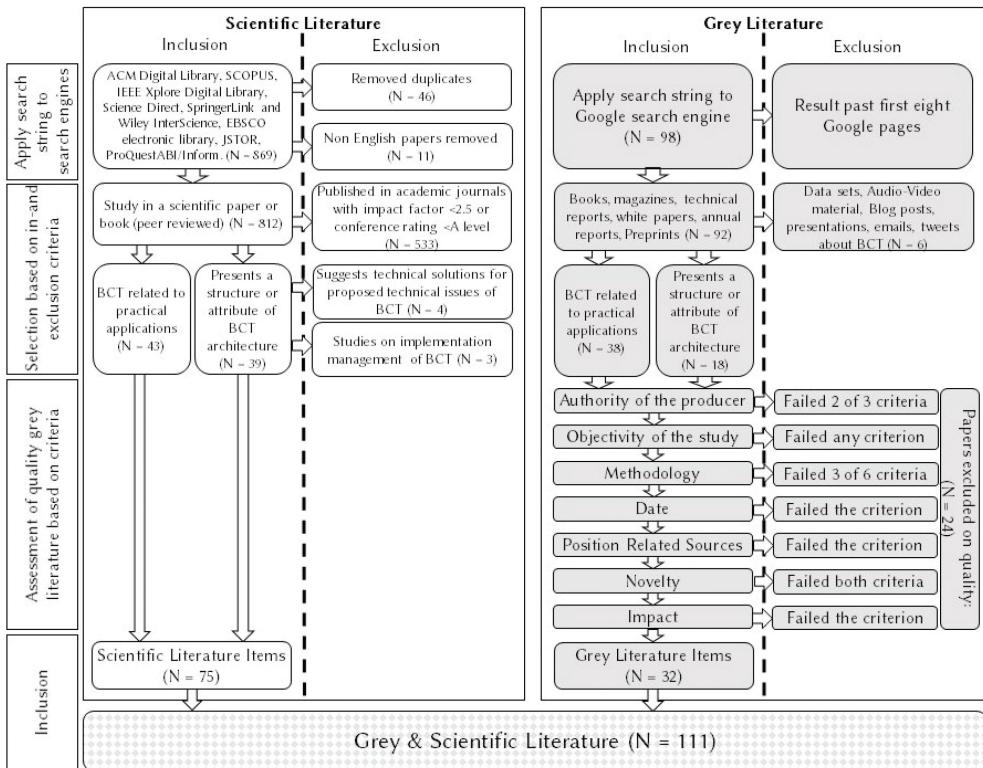


Fig. 2. Sample Search and selection strategy, a process model.

(1)“Blockchain”  $\vee$  “Blockchains”  $\vee$  “Distributed”  $\vee$  “Decentralized” (2) “Ledger”  $\vee$  “Technology”  $\vee$  “Database” (3) “Applications”  $\vee$  “Use Case”  $\vee$  “Implementation”  $\vee$  “Example”  $\vee$  “Case Study” (4) “Architectural”  $\vee$  “Architecture”  $\vee$  “Form”  $\vee$  “Fabric”  $\vee$  “Structure” “System”  $\vee$  “Design” (5) “Choices”  $\vee$  “Options”  $\vee$  “Decisions”. When combined, the preceding terms were used in the following search string:

$$[(1 \wedge 2) \wedge 3 \vee (4 \wedge 5)] \quad (1)$$

In the next stage (2), the search string has been applied to the following scholarly search engines: ACM Digital Library, SCOPUS, IEEE Xplore Digital Library, Science Direct, SpringerLink and Wiley InterScience, EBSCO electronic library, JSTOR knowledge storage, and ProQuestABI/Inform throughout March in 2018. The final stage (3) of the systematic review, the initial results were screened against inclusion and exclusion criteria that are shown in Figure 2 (Selection based on in-and-exclusion criteria).<sup>2</sup> For brevity’s sake we have not included a more elaborate description and rationale behind these criteria in the paper, but they can be accessed online (see Appendix A for more details).

For the second part of the MLR, to identify all relevant Grey literature, we established another protocol to filter and extract the GL using the guidelines suggested by Garousi et al. [28]. The protocol has been conducted in three stages: (1) search process, (2) source selection, and (3) study quality assessment. The implementation of these steps can be found on the right-and side of Figure 2.

<sup>2</sup>The (N = followed by a number) in Figure 2. represents the number of papers included or excluded based on the selection criteria.

In the first stage, we applied the search string to the Google search engine. The search process yielded 8.330.000 results when applying the first search string (“Blockchain”  $\vee$  “Ledger”  $\vee$  “Applications”). Because of the significant amount of results we initially limited our review to the first eight pages (20 results per page) provided by the Google search engine. Incrementally the next pages thereafter have been reviewed using inclusion and exclusion criteria related to the type of grey literature source (e.g., books, magazines or video files), (see Figure 2). Thereafter the pages were incrementally reviewed by title and abstract, starting from the first results page using the inclusion and exclusion criteria depicted in Figure 2 (selection based on in-and exclusion criteria). If <50% of the results on a page were not relevant for this research, then the search was stopped there. We further refined the GL studies we obtained from the first eight Google pages using the same inclusion and exclusion criteria.

During the second phase, we assessed the quality and relevance of the sources of the primary GL we obtained, since it cannot be assumed that the quality of GL is guaranteed. Exclusion criteria suggested by Garousi et al. [28] have been used for this purpose (see Figure 2).

The exclusion criteria used consist of seven quality categories ranging from the authority of the producer to the objectivity of the study that can be found in Figure 2 under assessment of quality grey literature. Combined these quality categories encompass 17 criteria that have been assessed one by one for each GL item. An overview of all the quality categories, quality criteria, and how many of these criteria had to be satisfied to include the item can be found in Table 1.

The authors of this study (viz., the first two authors of this study) have indicated whether a GL item (a) satisfied or (b) did not satisfy a criterion. In the cases where one of the criteria could not be assessed (e.g., because this information was missing) we have assessed these criteria as if they did not satisfy the criterion. GL items that did not satisfy the threshold for each quality category were excluded from the sample. After the selection process we merged the grey-and scientific literature into one sample as the literature under review.

### 3.3 Data Analysis

This section details the analysis methods enacted to address our research questions.

**3.3.1 Formal Concept Analysis.** To address RQ1, a Formal Concept Analysis (FCA) approach was adopted. FCA is a systematic approach to derive a formal ontology or concept hierarchy from a set of objects and their attributes [82]. A complete description of the FCA method employed to attain our definition of BCT can be found in Appendix B.

**3.3.2 Grounded-Theory Analysis.** For steps 2 and 3 of the MLR, and to address RQs 2, 3, and 4, a Straussian Grounded-Theory (GT) approach [30] was adopted. In the scope of Straussian GT, a series of systematic steps are enacted to allow a theory to emerge from the data (hence, “grounded”) using codes. For our research each code represents a concept or theme related to BCT. Whenever a paragraph in the literature under review represented one of these concepts or themes, the related appropriate code has been attached. In GT this process known as “coding” and includes the phases that are described in Appendix C.

### 3.4 Inter-Rater Reliability Assessment

We employed Krippendorff coefficient (or  $K-\alpha$ ) [33] to evaluate the inter-rater reliability of the inclusion and exclusion of SL items, the in- and exclusion of GL items, the quality assessment of the GL, and the coding process of the pilot study. The coefficient measures the agreement between two ordered lists of codes that have been applied as part of content analysis. The methods used to assess the inter-rater reliability between the raters, and the results thereof can be found in Appendix D.

Table 1. Quality Criteria Grey Literature

Category	Exclusion Criteria	Criteria to Satisfy
Authority of the producer	The publishing organization is reputable, or the individual author is associated with a reputable organization	2/3
Objectivity of the study	The author has published other work in the field The author has expertise in the area (e.g., job title)	3/3
Methodology	The statement of the sources is objective  There are no vested interests Conclusions are supported by data	4/6
Date	The source has a clearly stated aim The source has a clearly stated methodology The source is supported by authoritative, documented references Limits are clearly stated The work covers a specific question The work refers to a particular population	1/1
Position related sources	The item has a clearly stated date Key related GL or formal sources have been linked/discussed	1/1
Novelty	The item enriches or adds something unique to the research The item strengthens or refutes a current position	1/2
Impact	The GL source should have citations and backlinks to substantiate the arguments made in the study	1/1

### 3.5 Sample Selection Results

This section outlines the sample results of our search strategy (Section 3.2). First, the section presents the distribution of the sample between grey and scientific literature, along with the distribution the publication venues per year. Thereafter, the section showcases a frequency analysis of the topics discussed in the papers under review along with an overview of these trends per year. Finally, the distribution of these topics between grey and scientific items is presented and discussed.

**3.5.1 Publication Venues and Distribution Literature.** After the assessment of the GL, the search strategy yielded a total of 33 GL items and 78 of scientific peer-reviewed papers selected for this study. From this point onward, we no longer distinguish the results whether they were derived from GL or SL but flesh out results over the total of 111 studies be the object of this research. Figure 3 depicts the number of SL and GL per publication year. The results depicted for the year 2018 have only been collected until March 2018 and therefore might skew the results.

Figure 3 shows that from 2008 until 2013 BCT has gained little attention from either practitioners or scholars. The figure show an overall increase in the SL and GL published from 2014 onward. Furthermore, the statistics show that there is a growing interest from the scientific community for research in BCT. More specifically, from 2017 onward twice as many articles have been published

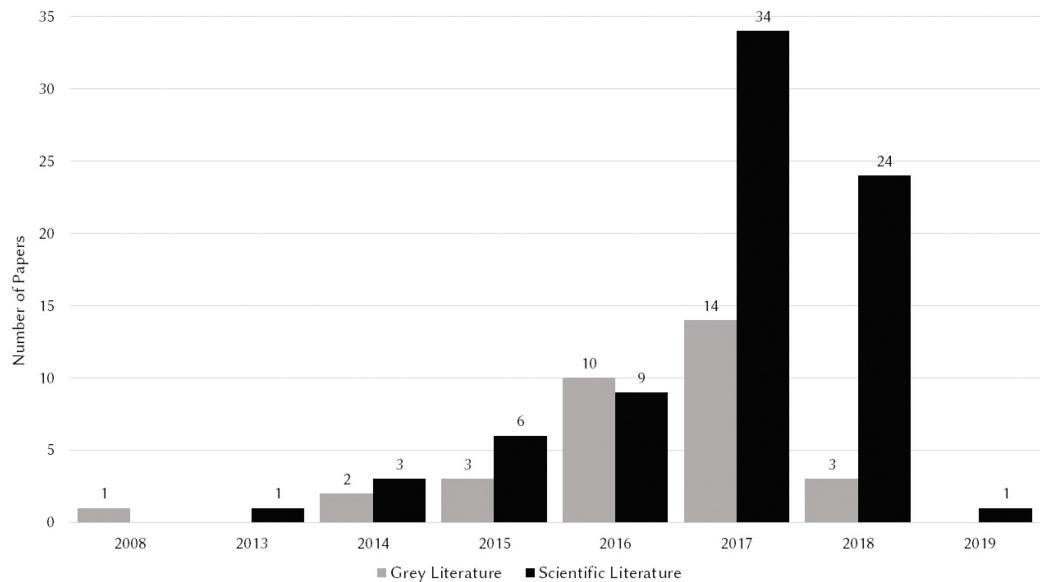


Fig. 3. Sample results: Grey and scientific literature across primary studies.

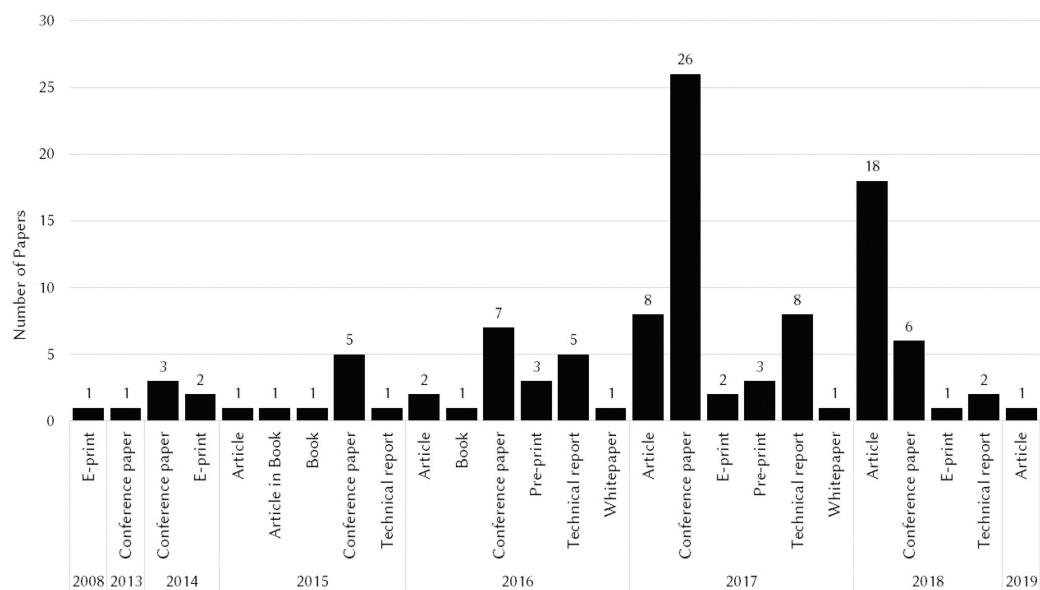


Fig. 4. Sample results: publication venues per year.

as compared to the years before. However, these results also indicate that research in the field of BCT is still in its infancy given that from 2008 to 2016 little scientific work has been published. The search and selection results also indicate an increase in the amount of practitioners literature being published. Furthermore, the sources from which these research items were identified for this study are diverse (see Figure 4), ranging from articles in technical magazines, books, and technical

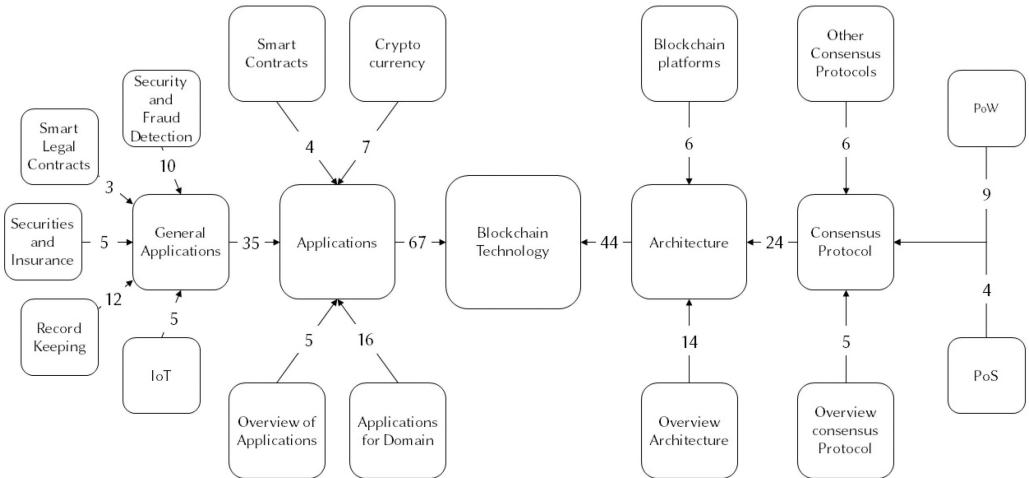


Fig. 5. Sample results: topics frequency analysis.

reports alike. The majority of items, however, were published in conference proceedings and reflect white literature.

In the years directly following 2008, i.e., the introduction of BCT, publications on the topic were almost evenly distributed among different sources (e.g., books, conference papers). However, as of 2016, books on BCT have not been found by this research. Although the publication of conference papers shows an increase in 2017, articles published in scientific venues are more gradually increasing in frequency.

**3.5.2 Topics in the Field of BCT.** Figure 5 shows the the main topics of the papers under review, as elicited using a Grounded-Theory approach. More specifically, the blocks in Figure 5 represent the topics found in the literature, while the number on the arrows between the blocks represents the *weight* of the topic, in terms of number of papers where those topics were coded. The direction of the arrow itself depicts under which of the composed main topics the sub topics are categorized. Among the items selected for this study, BCT-based applications are strongly represented (see left-hand side of Figure 5) while items on BCT architecture and are represented slightly less, with a ratio of 2/3.

In terms of applications, five sub-themes can be distilled: (1) cryptocurrencies, (2) smart contracts, (3) papers that provide an overview of BCT, (4) literature that suggests applications of BCT for specific domains, and, finally, (5) literature that presents general applications for BCT. Interestingly, Figure 5 again shows that the applications of BCT for a specific domain and general applications of BCT has gained considerable attention. On the one hand, literature on BCT for specific domains encompasses complex domains such as e-government, the financial sector, and relief development [17], while literature on the general topic of smart contracts is rather limited. However, the specific technical architecture literature over BCT reflects five categories: (1) security and fraud detection, (2) smart legal contracts, (3) securities and insurance, (4) record-keeping, and (5) the Internet of Things (IoT). Among these categories the majority of papers has been published on utilizing BCT for record-keeping, that is, registering certain data on the blockchain to ensure its immutability. Another major focus of research on BCT applications is security and fraud detection, and these items focus on using blockchain for safe distribution of data among peers. A

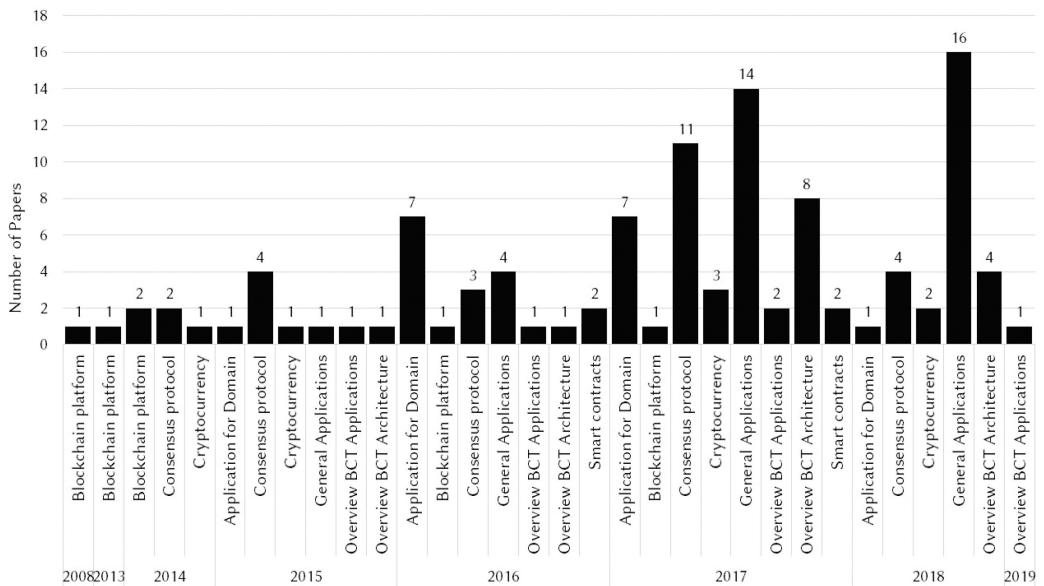


Fig. 6. Trends in BCT publication topics.

more detailed description of the contents of literature of applications for BCT can be found in Section 6.

**3.5.3 Trends in Publications on BCT.** Four key trends in the literature under review can be observed (see Figure 6). First, there is a balance in the distribution among topics even though the overall number of works on BCT is steadily increasing.

Second, two exceptions are (a) works related to applications for specific domains and (b) general applications research, as previously discussed—for these, the years 2016 to 2018 have seen a tremendous increase in research and practical work.

Third, publications on blockchain-based smart contracts have seen a rise only from 2016 onward. A similar observation can be made for items related to cryptocurrency as only one paper was published in 2015, with three papers being published in 2018.

Last, since 2014 an increasing amount of papers on the topic of consensus protocols have been published, a trend that continues to date.

What is more, in terms of past publications by both practitioners and scholars have predominantly been focused on applications for BCT (see Figure 7) with more of the grey literature being published on applications (72%) as compared to scientific literature (55%).

With respect to engineering research focus, scholars have focused their efforts on consensus protocols, whereas practitioners have mostly presented works on blockchain platforms or overviews of BCT architecture.

#### 4 A SYSTEMATIC DEFINITION OF BLOCKCHAIN TECHNOLOGY

So far we have referred to BCT as a single technology; however, BCT is a clever combination of several technologies and elements and there is no consensus on the definition of a blockchain [35], with a precise definition of blockchain technology often subject to controversial and subjective opinion [76]. Stemming from the data available to us, we strived to construct a rigorous definition of BCT based on (a) the identified software elements drawn from literature, (b) relations among

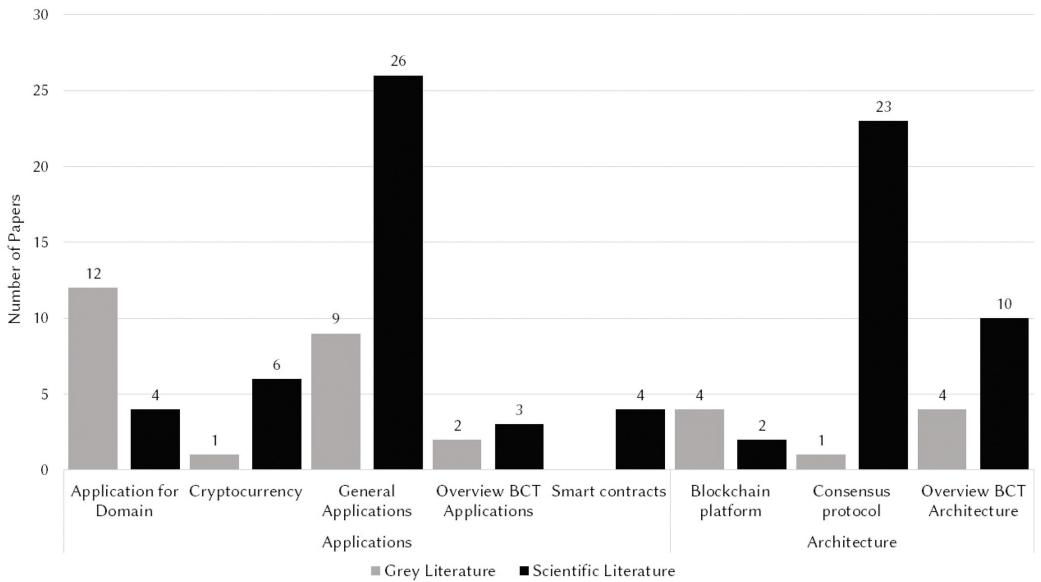


Fig. 7. Distribution topics scientific and grey literature.

them, and (c) their properties<sup>3</sup> reflecting the software architecture research and practice state of the art [8].

First, our data indicate that *Blockchain technology* is a form of *distributed ledger technology*, deployed on a *peer-to-peer network* where all data are replicated, shared, and synchronously spread across multiple peers. The technology allows actors participating in the network to perform, sign, and announce *transactions* by employing *public key cryptography*. Transactions are executed following a *consensus protocol* operated by specific nodes to ensure the validity of transactions requested by other peers in the network, and to synchronize all shared copies of the distributed ledger. During a consensus protocol execution, the data of valid transactions, along with other required metadata concerning the network, and the hash of the previous block are bundled into a *block* using *hashing functions*. The essential and key property reflecting BCT architectures is that each block contains the hash of their predecessor, therefore linking all prior transactions to newly appended transactions; the blocks therefore form a *chain* with the aim of establishing a tamper-proof historical record.

## 5 BLOCKCHAIN TECHNOLOGY: ARCHITECTURE ELEMENTS

This study examines the blockchain architecture landscape, arranging the elements found in the literature through the well-known 4+1 software architecture framework introduced by Kruchten [47]. The Kruchten framework delineates the comprehensive interplay of relations, properties and software elements in BCT and encompasses five views, namely (1) logical view, (2) development view, (3) process view, (4) physical view, and (5) a use case view. In Appendix E, a more elaborate description of these views is provided.

<sup>3</sup>Public-key cryptography is commonly used to describe the exchange of information using a set of private and public keys. Hence, when constructing a definition the term “public-key cryptography” has been used instead of private and public keys as an attribute of BCT.

Using the logical view, we first delineate the architectural elements to present the functionalities that various end users ultimately use from a blockchain. Further on, the development view describes how building BCT can be divided into smaller chunks of programmable code. Subsequently, the process view shows how IT systems behave during run time and is of interest to system integrators that need to know about the thread of control to execute operations utilizing BCT. Beyond that, the physical view is of interest to system engineers that maintain overall blockchain system, also, given that BCT completely resides on a P2P network its different arrangements is discussed in the physical view. The use-case view of the 4+1 model is recapped later in Section 6.

### 5.1 Logical View

The logical view emphasizes on the functional requirements and services the system should provide to its end users [47]. Decomposition of the architecture aids in identifying the elements that are common across the system. We used an ontology for BCT as proposed in References [23, 31] to organize the discussion of its main elements. A more elaborate, and in-depth description of BCT elements is provided online (see Appendix F how to access the material).

Blockchains are transaction oriented. Transactions in a blockchain system are executed using *public key cryptography*. *Cryptographic Hash functions* are used for the purpose of many operations, such as signing transactions (SHA-256 in the Bitcoin case [66]). The peers in the P2P network, also referred to as nodes are devices capable of processing and verifying transactions. Depending on the permissions all nodes or a specific subset of nodes validate transactions. The *permissions*. There exists at least three categories of blockchain networks [89, 93]; Public, private, and consortium networks that have different arrangements in terms of their permissions. On a blockchain transactions are stored in *blocks*. Each block is linked to its predecessor known as *parent block* by including its blockheader hash to form an integral chain of blocks that can be traced back to the first, or *genesis block*. Hence the term “blockchain” technology. Novel blocks are generated using a consensus protocol. Provided that the transactions included in the newly proposed block are valid, each new block enhances the security guarantees of the block before it [64, 66, 67]. Updates and changes to the software of a blockchain are called *forks*.

### 5.2 Development View

Existing blockchain networks can be leveraged to build *decentralized applications* (DApps) upon that use their services. Developers seeking to build their own blockchain platform have to program multiple software packages. Enabling transactions forms the basis for any blockchain network. A wallet needs to be programmed to allow clients of the platform to interact with other peers in the network. An *address propagation* method should be installed for nodes to interact. Next the nodes need to connect via *peer discovery*. Another aspect is the mechanism for *propagating transactions* [10].

Data with regard to transactions can be stored in two ways: As a first method, like Bitcoin, one can choose to add data into transactions. Another second method is to add data into contract storage like Ethereum [88]. Finally, an existing consensus protocol can be selected to process transactions or the protocol can be designed from scratch. Online material that can be found via a link in Appendix F further describes the development view.

### 5.3 Process View

The process view specifies which thread of control execute the operations of the classes identified in the logical view. The consensus protocol is at the heart of all BCT processes, since it allows for the enactment of transactions and ensures that the distributed ledger remains consistent. Online

materials to whom a link is provided in Appendix F delineates the steps, issues, and potential variants of the consensus protocols discussed in this section more in detail.

**5.3.1 Practical Byzantine Fault Tolerance (PBFT).** PBFT is mostly used in a private setting for permissioned blockchains, because it assumes authenticated nodes [25, 89, 93]. The protocol itself is exclusively based on communication, and nodes go engage in multiple rounds of communication to reach consensus [25]. Nodes do not get a reward for achieving consensus, rather in the event of malicious behavior by an authenticated node it can be held legally accountable [35, 73]. A primary leader node mines the blocks. The leader can be changed by other nodes via a “view-change” voting protocol, in the occurrence of a crash or when it exhibits malicious behavior [18, 64].

**5.3.2 Proof-of-Work.** (PoW) is often referred to as the *Nakamoto consensus protocol* [37, 54, 56, 70]. The PoW consensus protocol is designed for the case where there is little to no trust amongst users of the system [90]. Public blockchains need to have a high degree of Byzantine fault tolerance as users cannot trust one another.

Consensus in PoW is achieved through a hashing competition between miners. Competing miners need to commit computing power to calculate the solution to the same mathematical problem. To incentivize miners to participate in the consensus process the miner that is the first to find the solution to the mathematical problem reserves the right to publish the next block, and is rewarded by an amount of cryptocurrency [11, 65, 89, 90]. In addition, the miner to win the competition with its peers is also be able to collect the transactions fees that were paid by clients.

Finding the solution to a PoW problem is a computationally arduous process for which there are no shortcuts [67, 90]. The solution to the problem is hard to find, yet easy to check once they have been found [56]. Given that only one miner can win the competition and is rewarded the other nodes have simply wasted resources (CPU power and energy) in their attempt [29, 64, 77, 89, 90]. In addition, because the difficulty of PoW problems increases over time makes it even harder to win the competition [90].

**5.3.3 Proof-of-Elapsed Time (PoET).** PoET is designed to address the inefficiency of PoW and replaces it with a protocol that is based on *trusted hardware*. A node that uses trusted hardware, however, can be checked for certain properties such as whether it is running a certain software. This aids in relaxing the trust model in settings were the Byzantine’s Generals Problem might be present [25]. Sawtooth Lake, a project by Hyperledger, leverages Intel’s *Software Guard Extensions* (SGX) to establish a validation lottery that makes use of their CPUs capability to render a timestamp that is cryptographically singed by the hardware [36].

**5.3.4 Proof-of-Stake (PoS).** As a response to the limitations of PoW the BCT community has turned toward Proof-of-Stake (PoS). The PoS consensus protocol has been introduced for public settings [64] with the aim to safeguard against Sybil attacks and malicious behavior by untrusted nodes [25]. The PoS protocol offers a more efficient and environmental friendly alternative to PoW as computing power is partially substituted by virtual resources (e.g., cryptocurrencies) that miners must invest to propose blocks [29, 52, 77, 92]. Rather than using computer power as a scarce resource to generate security, Proof of Stake uses the scarcity of the coin itself. Therefore nodes that participate in a PoS consensus protocol are more commonly referred to as *forgers* instead of miners [6, 53].

The idea behind the PoS model is that the more assets (e.g., cryptocurrency), or *stake* a node has, its incentive to undermine the system diminishes, because subverting the system would inherently mean that the worth of the nodes’ stake would decrease [61, 90]. Logically, this implies that one cannot participate in the consensus protocol without owning a stake [29]. A shared commonality of all PoS variants is that nodes that have more stake have a higher chance of generating new

blocks [52, 64, 90, 93]. In other words, the more skin a forger puts in the game the higher its reward will be.

**5.3.5 Delegated-Proof-of-Stake (DPoS).** Delegated Proof-of-Stake introduces another variant of PoS [64, 93]. In DPoS, stakeholders elect delegates, referred to as *witnesses* to forge and validate blocks in round-robin fashion [53].

Compared to PoW and Pos, DPoS is more energy efficient. Further, because the voting about the validity of a block is delegated and fewer nodes are needed to validate, the blocks can be confirmed more quickly. Hence, as compared to PoW and PoS, DPoS has a low latency. Moreover, parameters including block size and block intervals can be adjusted by *committee members* of the governance board. When a delegate acts malicious this dishonest delegate can be voted out by all the other nodes [53, 93].

**5.3.6 Zero-Knowledge-Proofs.** Recently, different Zero-Knowledge-Proofs – (ZKP’s) based BCT networks have been proposed to preserve users’ anonymity and confidentiality of transactions [89]. In general, ZKP’s aim to confirm a statement about a transaction such as “This is a valid transaction” without revealing anything about the transfer (statement) itself or the parties involved [35, 89, 92]. Zerocoin was the first initiative with the aim of providing transaction unlinkability using ZKP’s [25]. Similarly to Bitcoin, Zerocoin uses the PoW consensus protocol to validate transactions. A *cryptographic mixer* is implemented for Zerocoin to conceal the links between a zerocoin and the corresponding Bitcoin.

Building on the ZKP approach as a foundation, Zcash extends the privacy guarantees and improves the efficiency (throughput and latency) of Zerocoin. Zcash uses a variant of the PoW called Equihash. Transactions made using Zcash, including the split and merge transactions, are fully private [25]. Zcash employs a technique called *Zero-Knowledge-Succinct Non-Interactive Argument of Knowledge* (zk-SNARKS) to provide these privacy guarantees [64, 92] that are a specific type of ZKP.

## 5.4 Physical View

The physical view is concerned with the topology of software components and their physical connections. Electronic devices known as nodes constitute to a blockchains’ P2P network and are the only physical connection to the non-digital world. P2P networks on which blockchain platforms are run have different arrangements; First, the network can be categorized on the basis of permissions (authorization). Second, networks can be categorized with regard to their accessibility. Permissions to perform operations on the blockchain might differ ranging from allowing anyone to read, write and to partake in the consensus protocol to only one of these permissions. Control over these permissions can be confined to a distinct group of nodes or all nodes.

As the name suggest *Permissionless* grant permission to all nodes in the P2P network to read and write transactions. *Permissioned* blockchain platforms have confined and idiosyncratic permissions for their nodes [63, 75, 89, 90].

The P2P network can also be described from the perspective of network accessibility. In the literature three categories of P2P networks can be distinguished that are coupled to a permission model [13, 23, 35, 59, 89, 93]. A *Public blockchain*, like Bitcoin or Ethereum have *open network access* meaning that anyone willing is allowed to join the network. *Private blockchains* are blockchains networks that are owned by one organization. Contrary to public blockchains access is confined. *Consortium blockchains* are similar to private blockchains in the sense that nodes first need to be authenticated before granted access to the network. However, consortium blockchains allow nodes from different organizations to access the blockchain network [64, 74, 89]. A more elaborate

description of both models and blockchain networks can be found in online materials accessible via a link in Appendix F.

## 6 BLOCKCHAIN USE-CASE VIEW: MAIN USAGE SCENARIOS

The use-case view aims at providing a description of an architecture by illustrating an essential set of use cases and scenarios for their usage. Our data suggest there are mainly three flavours of BCT applications: (1) cryptocurrencies, (2) smart contracts, and (3), general-purpose applications. Following these versions we further categorize BCT applications [77]. For instance, the general-purpose applications of BCT can be arranged into five additional categories that encompass: (1) Security and Fraud Detection, (2) Securities and Insurance, (3) Record-Keeping, (4) Internet-of-Things, and (5) Smart Legal Contracts. A complete catalogue of BCT applications can be found in Appendix F.

## 7 BLOCKCHAIN TECHNOLOGY: MAIN ARCHITECTURE PROPERTIES

The scope of our analysis revealed eight essential architectural properties with a directed mutual influence relation evident from the state of the art. In fact, stemming from the relations our GT analysis we marked with a  $\rightsquigarrow$  operator the mutual implication relation evident between the following couples of properties:

- *Decentralization*  $\rightsquigarrow$  *Disintermediation*. In traditional centralized transaction systems each transaction needs to be validated by a (trusted) third party (e.g., a bank). The decentralized workings of BCT enables the direct transfers of digital assets between two counter parties without this third party leading to direct disintermediation [29, 67, 75, 93].
- *Programmability*  $\rightsquigarrow$  *Automation*. BCT allows for the execution of pre-defined conditions that are automatically executed once certain conditions have been met. BCT enabled smart contracts extend this concept further by allowing (Turing complete) programmability of transactions [19, 36, 75]. The Bitcoin blockchain predominantly offers a service to exchange cryptocurrency and, accordingly provides limited support for smart contracts Blockchains like Ethereum or Kadena offer a fully programmable smart contract environment [35, 61, 67, 75]. Offering smart contracts as a service, however, adds another layer of complexity; smart contract execution puts a higher strain on the data storage requirements, throughput and latency of a blockchain network [21, 23, 35]. Furthermore, arbitrary code leaves room for human errors, and thus increases the chances of bugs [23, 29, 54]. In sum, the degree of automation depends on the services provided, which is closely linked to the *design* of a blockchain platform but at the cost of additional complexity.
- *Transparency*  $\rightsquigarrow$  *Auditability*. Each node in a blockchain P2P network holds a complete copy of the distributed ledger making all transactions transparent [11, 29, 77, 87]. However, for permissioned blockchains permissions to read the ledger can be confined to increase transaction privacy. Decreasing the transparency of the transaction records makes permissioned blockchains less auditable [59, 87, 89]. In short, the auditability of the network depends on the permission *arrangement* of the P2P network.
- *Immutability*  $\rightsquigarrow$  *Verifiability*. The entire history of transactions performed is recorded and stored in blocks. Given that these blocks are cryptographically chained using hashes, the record becomes immutable [59, 67, 75, 76, 79]. Provided that the entire history of transactions is auditable, the proof that any transaction has (not) taken place in the past is thus verifiable, since blockchains are append only [29, 76, 81]. An insecure consensus protocol that allows for the introduction of blocks containing double spend transactions could jeopardize the immutability of the ledger. Therefore the immutability of a blockchains

distributed ledger depends on how transactions are *processed* during the consensus protocol [67, 75, 90].

## 8 BLOCKCHAIN TECHNOLOGY: CHALLENGES AND OUTLOOK

Despite being a promising novel technology, currently BCT faces several challenges that inhibit widespread adoption. This section highlights and discusses the challenges evident from the literature.

### 8.1 Latency

One of the challenges BCT faces is that most consensus protocols have a high latency, meaning that the time between the submission of transactions and their confirmation is high [29, 31, 35, 61, 65, 89]. This is due to the fixed blocktime interval for most blockchain networks. Effectively this means that on average it takes the Bitcoin network roughly 60 minutes before transactions are settled and can be regarded as final [35, 89]. Ethereum has made significant progress in this area using the Greedy Heaviest Observed Subtree (GHOST) protocol by increasing the block interval to 14 s and transaction finality after 12 blocks [29, 89].

What is more, currently the finality for clearing and settling transactions is a legally defined moment. When using BCT to enact transactions settlement finality is probabilistic; the longer a transaction is considered settled by network participants, the less likely it will become that the transaction will be reversed or declared invalid [35, 63]. Clearly, these two arrangements are at odds. A direction that is currently being explored to concurrently address the throughput and latency issues of BCT is that of *sharding* the mining network. ELASTICO [55] is an example of a consensus protocol that shards the mining network. When sharding the network miners are uniformly partitioned into smaller committees that process a specific set of transactions. Accordingly transactions can be processed in parallel and thus throughput capacity can be increased.

### 8.2 Throughput

The maximum throughput of transaction has also been shown to be a challenge [35, 64, 87, 89, 93]. The concurrent throughput challenges for BCT are closely related to those of the latency. At the time of writing the Bitcoin network can reach a throughput of seven transactions per second [25, 35].

Yet again this problem is related to the blocktime interval but also to *blocksize*. The size of a block determines how many transactions can be included. For Bitcoin the size limit of a block is 1 MB [89]. Recently, there have been proposals to increase the throughput of the Bitcoin blockchain by increasing the blocksize from 1 to 8 MB [89]. Proponents and opponents of this proposal have interchanged various arguments that so far has reached no conclusive upper-hand [79]. By implementing the *GHOST protocol* Ethereum has managed to improve its throughput capacity to 15 transactions per second, because the block time interval is smaller (14 s). Rather than following the longest chain, in GHOST a miner weights the branches in terms of the computational power spent to create them and chooses the better one to follow. Another promising novel development is the introduction of *off-chain payment channels* such as Raiden,<sup>4</sup> Bitcoin Lightning,<sup>5</sup> and Sprites [62] that enables two parties to directly and privately maintain a two-party micro payment channel. Khalil and Gervais [40] extend the concept of off-chain payment channels by suggesting a novel approach that enables the refunding of existing payment channels when they are depleted without performing a transaction on the blockchain network. Recently the segregated witness

---

<sup>4</sup>[www.raiden.network](http://www.raiden.network).

<sup>5</sup>[www.lightning.network](http://www.lightning.network).

(SegWit) proposal has been suggested in the Bitcoin community to change the internal design of blocks to increase the throughput of transactions. The proposal entails separating (segregate) signatures (witnesses) from the remainder transaction data. In this manner, the size of the witnesses does not add to the data size limit of the blocks [89].

### 8.3 Data Storage

Another challenge that is pointed out by both practitioners and scholars alike is how to cope with the evergrowing need for *data storage space* [29, 35, 50, 77, 94]. This challenge mainly stems from the fact that to verify transactions, a node needs to be aware of the whole blockchains' history. If Bitcoin were to process an equal number of transactions as Visa, then the amount of storage required would grow by 214 PB per year [77]. Some suggestions for improvement of data storage have been made such as the introduction of *lightweight clients* that do not download the complete record of transactions. Instead, lightweight clients download only the blockheaders to validate transactions. To verify transactions these nodes use a technique called Simplified Payment Verification (SPV) [90, 93].

### 8.4 Data Privacy

Preserving privacy of participants and confidentiality of their data has turned out to be a fundamental challenge [22, 35, 60, 65, 76, 87, 89, 93]. Although transparency is one of the key characteristics of especially public blockchain networks it is at odds with privacy. For public blockchains by design every transaction needs to be visible to every participant for the sake of public verifiability [3, 35, 90], though they can be encrypted and the identity of the user is hidden. To address this problem private and consortium blockchains such as Hyperledger [3] and Corda [12] with a permissioned model have been introduced [67]. Another approach solve this problem is the usage of mixers and ZKP (see Section 5.3.6) [5]. In a study [35] 57% of the respondents stated that implementing privacy-enhancing techniques in their BCT systems is planned for the future. Out of these respondents 78% have expressed the desire to implement zero-knowledge proofs (ZKP). A second privacy challenge is that a blockchain ledger is *immutable*; Once a transaction has been stored in a block it cannot be removed. Further, in permissionless blockchain every node is able to view all transactions and, consequently explore the entire history of transactions. The General Data Protection Regulation (GDPR), however, enforces restrictions on how information about EU citizens may be used and stored [68]. One of the rules that would be difficult to comply with is the "right to be forgotten" that allows an individual to demand the erasure of information under certain conditions. Clearly, the immutability of a blockchains ledger is incongruent with the right to be forgotten [36, 76].

### 8.5 Governance

The governance of a blockchain with regards to *updating* its fundamental rules is problematic [67, 90, 92, 94]. A prime example is the ongoing debate within the Bitcoin community about the block size, which has ended in a stalemate [59, 79]. Even for centralized systems updating software can be difficult let alone when a system has many users, geographically dispersed, as can be the case with BCT [90]. Another classic example of the governance problems blockchain currently faces is the response of the Ethereum community to the DAO hack.<sup>6</sup> Due to unintended flaws in the semantics of a contract an attacker was able to obfuscate a large amount of Ether worth an

<sup>6</sup>The term hack must be qualified; The attacker exploited a vulnerability in the smart contract that allowed a split function (enabling the withdrawal of funds from the contract) to be called repeatedly to withdrawal more funds than entitled to. For further reading about the DAO hack the author recommend reading Annex B in Reference [36].

estimated \$50 million.<sup>7</sup> In response to the attack, a hard fork was proposed to recover the Ether, to which 89% of Ether-holding voters gave their consent. Some of the remainder non-consenting voters rejected this fork of the blockchain mostly for philosophical reasons, including the principle that a blockchain is immutable. These voters decided to use the unforked Ethereum blockchain resulting in a split into two separate currencies: Ether (containing the hard fork) and Ethereum Classic (no hard fork) [67, 90]. Another governance issue that needs to be addressed is that of *key management*; BCT is decentralized and as such when a user forgets their private key there is no central authority to recover it [36]. As a solution to this problem He et al. [34] present a wallet-management system based on semi-trusted social networks to recover wallets and the keys they hold. However, the true Achilles heel with regard to private key management is related to the wallets that store these key's; if the hardware on which a users wallet gets lost, targeted with malware, or is attacked, then the private key might get lost or stolen [5].

## 8.6 Usability

A more practical challenge that hampers the widespread adoption of BCT is the current lack of end-user support (BCT is hard to use and to understand) and adequate developer support (few developers tools available) [60, 65]. In line with these observations, research by Tapscott and Tapscott [79] indicates that many Dapps are not accessible to the average person and that interfaces are user-unfriendly. Further, in their study they suggest that there are approximately between 1,000 and 2,000 developers that understand how to develop Dapps. However, one of their interviewees stated that this number could perhaps increase by establishing creative educational programmes.

## 9 DISCUSSION

Our GT-based analysis was used to populate the illustrated 4+1 views, properties, and challenges of BCT. First, this section grounds the insights on the 4+1 views of BCT architecture elements, properties, and challenges through discussion. More specifically, the GT-driven 4+1 perspectives on BCT are deepened by discussing coding frequencies and trends of the concepts in literature. Second, the section presents observations we made in the scope of our analysis derived from examining the distribution of the topics found in the sample and synthesizing their contents.

### 9.1 A Grounded-Theory of Blockchain Technology

The first step of in the data analysis procedure was to apply an open code each time the literature reflected a concept (see Section 3.3). A frequency analysis of the open codes (how often certain codes have been applied) unravels which concepts are deemed important. Blockchain challenges were mentioned most, followed by smart contracts, and consensus protocols. These results show that blockchain challenges are widely discussed in the papers under review and that the technology has not yet come to full fruition. The frequency analysis further showed that smart contracts are most frequently discussed, and can therefore be deemed a key concept for BCT future evolutions.

The third most often applied code is that of consensus protocols and among them PoW. These findings resonate with the number of papers on the topic of PoW consensus protocols (see Section 3.5). Taken together, these findings show that literature is still predominantly focused on the PoW consensus protocol, whereas several other consensus protocols nowadays exist.

Surprisingly, the term permissioned blockchain is more often mentioned than permissionless blockchain, which is not mentioned as one of the top 10 concepts. However, these results can also

---

<sup>7</sup>Tapscott and Tapscott [79] argue that the total worth of the obfuscated Ether was around \$70 million, whereas Gatteschi et al. [29] suggest that it was the equivalent of \$60 million.

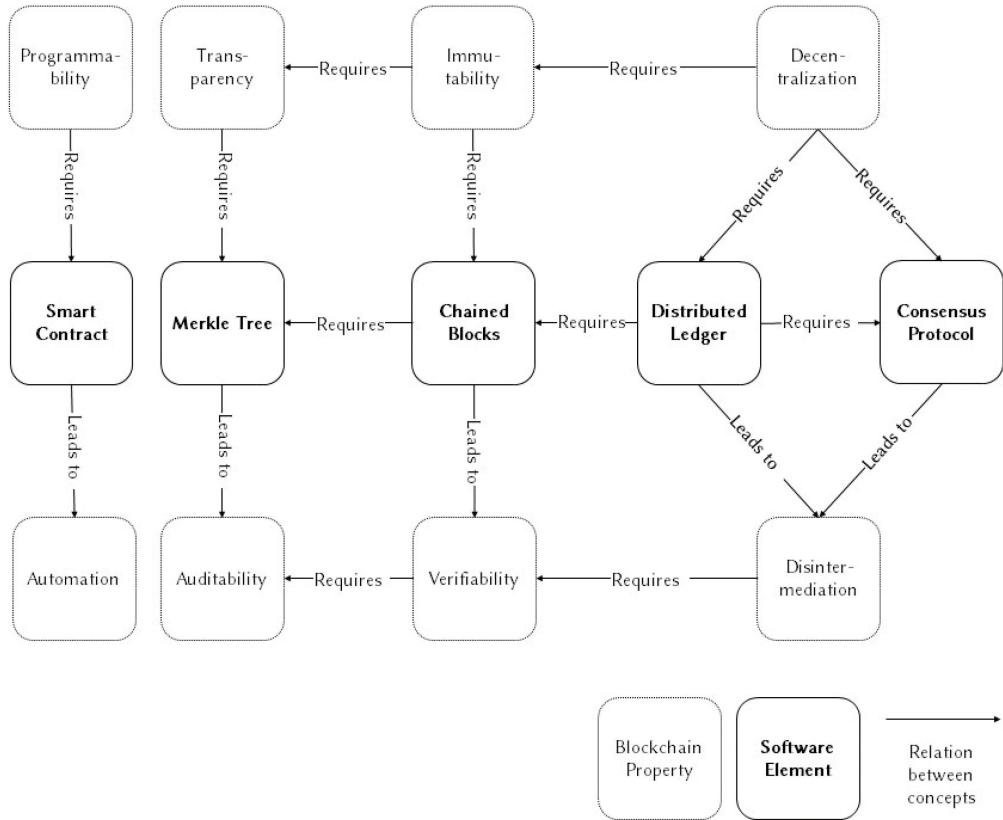


Fig. 8. BCT software elements and properties, an overview.

be attributed to the perception that blockchains generally are public and permissionless and as such that permissioned blockchains are an exemption to be specifically mentioned.

**9.1.1 Grounding the Logical View of BCT.** BCT encompasses several software elements that combined create the architectural properties of BCT. The results of the axial coding revealed what the intricate relation among software elements and properties.

The software elements of BCT work in concert to allow for secure transactions on a P2P network (see Section 5.1). We found that there is a strong dependency among these elements, which has been depicted in Figure 8 by the middle row of blocks: During a consensus protocol nodes verify transactions, which is not possible without the availability of a distributed ledger. In turn, the distributed ledger employs the concept of chained blocks that depends on a Merkle Tree to summarize the transactions. Because each software element is implemented to yield a particular property (see Section 7) these properties also are also connected. This relation is shown as the top row of blocks in Figure 8.

Decentralization requires a consensus protocol and distributed ledger but in addition immutability of the transaction records. The immutability of transaction records is dependent on the degree of transparency of a ledger (see Section 7). However, transparency does not only guaranty the immutability of the distributed ledger, but also leads to auditability of the transaction records (bottom row Figure 8). Verifiability of a blockchain requires auditability as without no proof can be provided that a transaction has not already been spent.

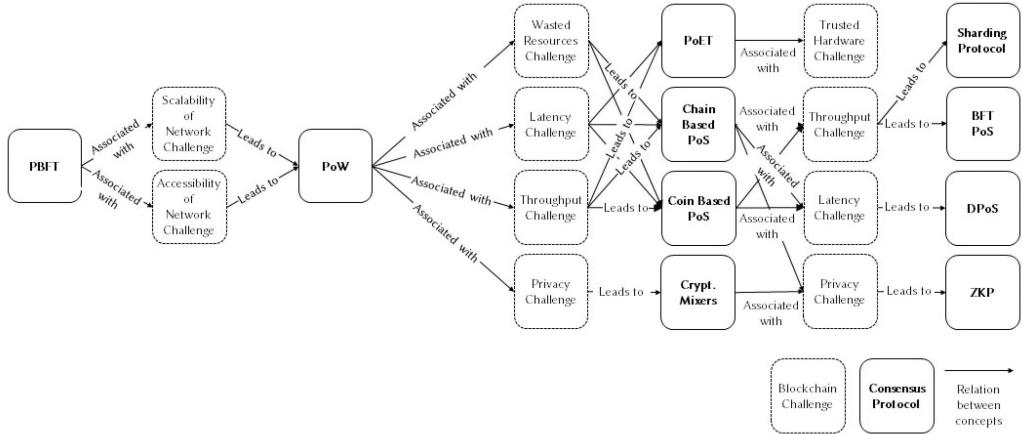


Fig. 9. Consensus protocols related to challenges in chronological order from left to right.

An interesting observation is that programmability does not seem to have a direct relation to the other properties of blockchain. Every blockchain has a certain degree of programmability [89]. However, arbitrary programmability of transactions requires the concept of smart contracts, which is optional for blockchain design [88]. The discussion of these results illustrates the complex relation among BCT software elements and properties. Considering these relations will be of importance to obtain the appropriate design when implementing BCT.

**9.1.2 Grounding the Process View of BCT.** Another relation that emerged from the axial coding process is the relation between BCT challenges and consensus protocols. In Figure 9, this relation is shown simultaneously with the introduction of consensus protocols in chronological order.

Our results from the axial coding process show that each consensus protocol developed after the PoW protocol aims to tackle a particular BCT challenge. Furthermore, unsatisfied with prior protocols to address the challenges of BCT novel consensus protocols were introduced over time. For instance, the development of cryptographic mixers clearly aimed at providing more transaction privacy. To facilitate fully private transactions, however, ZKP-based protocols were introduced. While both approaches increase or facilitate fully private transactions these protocols were not designed to achieve faster throughput as compared to PoW.

However, PoS-based protocols have aimed at improving throughput and latency performance as compared to PoW. Among the first of these attempts were the coin- and chain-based PoS variants (see Section 5.3.4). As the performance aspects of these PoS variants were still considered unsatisfactory BFT-based PoS, DPoS- and sharding-based protocols were introduced [53, 90]. None of the PoS, DPoS, or sharding consensus protocols address privacy of transactions.

Taken together the results show that there are have been desperate attempts to address BCT challenges. A uniform approach that simultaneously addresses specific sets of problems is still lacking. This implication might have some important ramifications for future research efforts, for example, several research efforts by practitioners and scholars have led to consensus protocols that only pursue to improve one favorable property while a unified agenda to develop a consensus protocol that is designed with both privacy and performance aspects (throughput and latency) in mind is still lacking.

**9.1.3 Grounding the Physical View of BCT.** The grounded view describes the design patterns of a blockchain network. As discussed in Section 5.3 the design of the P2P network that supports

the blockchain can differ. The frequency of the codes concerning BCT networks have been examined to establish a notion of their importance. What becomes clear from these statistics is that surprisingly permissioned blockchain networks are more mentioned than permissionless BCT networks. A closer examination of the papers in which the codes were used reveals that permissioned networks are predominantly employed to contrast the properties (negative and positive) of permissionless networks. Public networks are the third code that is most often applied, followed by private networks. These paradoxical networks are discussed in tandem to contrast their properties while Consortium networks are mentioned least. It must be noted, however, that the terms consortium and private networks are often discussed in the same context.

The trends on BCT networks also sheds an important light on the developments in the field of BCT. These trends have been constructed by counting each time a paper was coded the network type in a given year. Thereafter the result has been divided by the total number of papers in the sample for a particular year. We deemed this last step appropriate to account for the fact that over time more work on BCT has been published. Analysis of the trends regarding BCT networks codes (see online materials in Appendix C for more details) showed that in time more papers were mentioning private and consortium networks combined with permissioned models. When closer examining this trend we found that over time more papers were discussing private, consortium and permissioned networks. The literature under review [26, 29, 35, 74, 75] states that these network types were introduced as a response to the current challenges public oriented networks face (e.g., privacy and throughput). Therefore this trend can be explained by the fact that in time public blockchain challenges and limitations became more evident and thus alternatives were introduced.

**9.1.4 Grounding the Use-Case View of BCT.** The axial coding process of papers concerning BCT applications revealed that these applications have a distinct focus to which they have been categorized (see Figure 5). A large strand of literature is focused on cryptocurrencies and improving the interoperability between chains using different coins, or fluctuations in the prices of cryptocurrency. The focus of these papers lies on *transactions* carried out using a blockchain or between chains and what influences these transactions. For these papers therefore the cryptocurrency itself becomes the *specific* focus of study as they are the embodiment of blockchain transactions. This contrast with other papers that see cryptocurrency as a means to an end and tend to have a more general focus on employing the technology as a whole. Several papers suggest the use of BCT to enhance the utility of IoT devices. Here, BCT *primary* serves to improve, amongst things, security and connections of these devices. In turn, IoT devices could potentially be used for a wide array of other applications. Discerning whether BCT is the application or rather a means to an end helps to categorize papers in the field of BCT.

Another observation resulting from the axial coding process is that the general application categories found in the literature sample utilize different properties of BCT. Record keeping applications mostly benefit from the decentralized nature of BCT. Whereas applications related to security and fraud detection seem to employ BCT for the sake of verifiability. For securities and insurance applications the auditability of the transactions is most important. Thus, besides categorizing applications based on their primary focus, one should also take into account which of BCT's properties are predominantly utilized. This observation is important to take note of when developing BCT-based applications as some properties can be deemed more valuable for the design than others and can therefore traded-off against each other.

**9.1.5 Grounding the Properties and Challenges View of BCT.** BCT has been developed to attain certain properties yet also gave rise several challenges. We juxtaposed the results of a frequency analysis conducted using the codes applied for BCT properties and challenges to discuss their

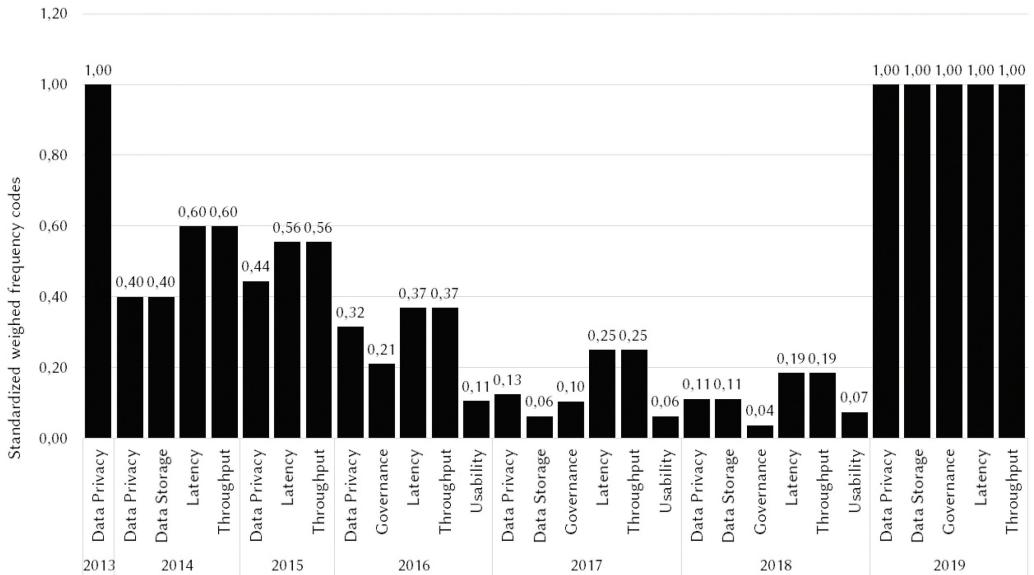


Fig. 10. Trends in BCT challenges from 2008 to 2019.

relation. The frequency of the codes related to BCT properties and challenges can be found in Appendix C.

Challenges related to privacy on a blockchain is mentioned most of all codes related to challenges. The immutability is the most frequent applied code related to BCT properties. Paradoxically, whereas privacy the most discussed challenge for BCT, the transparency of the ledger is mentioned as the second most important property of BCT. The decentralized nature of BCT is the single least applied code. Automatic execution is the least recurring code of related to a property of BCT. However, papers that predominantly discuss properties from the perspective of Bitcoin do not mention smart contracts, which is related to the concept of automatic execution.

A further analysis of the relation between the throughput and latency codes applied reveals that the coding of these two concepts coincides, resulting in an almost equal number of times these codes are applied. Governance of blockchain networks as a challenge is almost as frequently discussed as the technical issues of BCT (throughput and latency). The data storage and usability codes are less frequently used in comparison to the other codes concerning BCT challenges. Of these two codes, usability is least often mentioned as a challenge. These results can be explained by the fact that most papers either focus on the applications for BCT or on architectural aspects without regarding the users perspective.

An analysis of the applied frequency of codes over time reveals the trends in BCT challenges. Again, the analysis has been conducted by counting each time a paper was coded a certain challenge in a given year. After obtaining the results these have been divided by the total number of papers for a particular year. The weighted number of codes per are challenge and year are depicted in Figure 10 to provide a fine-grained perspective. The figure shows an overall decline between 2008 and 2019 in the mentioning of all challenges.<sup>8</sup> What is interesting, however, is that when the

<sup>8</sup>For the years 2013 and 2019, only one paper has been included in the literature sample that imbalances the standardized weighted frequency of the codes for these years. Hence, in the scope of this analysis the results for 2013 and 2019 should be regarded as outliers.

codes are weighted (i.e., per year in relation to total number of papers) the results demonstrate that throughput and latency are the most important challenges to address. From publications in 2016 challenges related to BCT governance and usability first emerge. This shows that from 2016 onward blockchain became mainstream adopted as the usability of BCT (for non-programmers) for the first time were discussed in References [63, 83]. Moreover, in that same year governance emerged from the literature under review [1, 63, 73, 83] as a challenge. This literature discusses that novel updates to existing platforms were required due to several reasons yet that this turned out to be a challenge. In tandem, the governance issues from a legal perspective are discussed, which means that BCT is was no longer regarded as a novelty but mainstream and a technology that needed to adhere to legal standards (see Section 8.5).

## 9.2 Highlights and Observations

9.2.1 *Chaining Data Using Blocks.* From the FCA conducted for this research we found that what sets blockchain apart from other forms of distributed ledgers is that transactions are stored in blocks. On the one hand, transaction data are stored in blocks to ensure the integrity of the distributed ledger. However, storing data in blocks also has its drawbacks.

Blocks have fixed data sizes and can only contain a limited number of transactions. Transaction that do not fit in the current block that is created have to wait to be processed and included in the next block. Given that both blocktime and blocksize for most blockchain networks are fixed these two parameters determine throughput capacity of a blockchain network [67]. The SegWit proposal (see Section 8) might partially solve this problem by expanding the size of blocks. However, to increase the throughput of public blockchain protocols further future expansions of the blocksize will be required. What the consequences are for the security of the network remains unknown, however, because there has been no empirical investigation to test these configurations.

Nodes on a blockchain network need to keep a complete history of all transactions made on a blockchain network to validate them. Because the number of transactions that have been made on a blockchain is growing over time the data storage demands grow in parallel. Consequently, in the long term it will be unsustainable for every node to keep the entire history of transactions [77]. Conversely, if only few nodes would be able to meet the store demands, then it would defeat the purpose of decentralization. A similar trend can be observed with regards to Bitcoin mining. Initially nodes mined blocks individually. However, due to increased mining requirements (CPU power) they eventually started collaborating in mining pools and mining became more centralized as a result. To address this problem some blockchains are utilizing the concept of checkpoints [53, 59]. Whether this solution aides in securing the network from attacks remains unknown. Another proposed solution is the use of lightweight clients. However, at least some nodes need to be full weight clients that keep track of all transactions and thus still have to burden themselves with storing large amounts of data.

9.2.2 *Consensus in the Wild.* Besides the consensus protocols discussed in the results section (see Section 3.5) we reported more consensus protocols such as *Proof of Work or Knowledge* (PoWoKs) [6], *Proof of Vote* (PoV) [52], *Proof of Sequential Work* (PoSW) [20]. However, none of these protocols is supported by mature implementations, practical application, or empirical evidence of operational characteristics. For the sake of completeness, however, Table 2 offers an overview of all primary consensus protocols (without any derivation, e.g., PoSW with respect to PoW) and their claimed features in the respective literature. The table articulates every consensus protocol (ID in column 1) using (a) the trust-level required in the P2P network setting, (b) the scalability (i.e., whether the network can scale in terms of the number of nodes [38]), (c) the

Table 2. A Complete Reference over Consensus Protocols

<b>ID</b>	<b>Trust-level</b>	<b>Scalability (#Nodes)</b>	<b>Byzantine Fault-Tolerance</b>	<b>Throughput</b>	<b>Latency</b>	<b>Example</b>	<b>Source</b>
PBFT	High	Weak	<33.3% of faulty replicas	<2000	<10 s	Hyperledger Fabric v0.6 <sup>9</sup>	[13, 64, 93]
RAFT	High	Weak	<51% of faulty nodes	>10k	<10 s	Corda <sup>10</sup>	[13, 64]
PoW	Low	Strong	<51% of computing power	<100	>100 s	Bitcoin <sup>11</sup>	[64, 93]
PoET	Low	—	<51% of computing power	—	—	Hyperledger Sawtooth <sup>12</sup>	[13]
PoS	Low	Strong	<33% of stake	<1000	<100 s	Tendermint <sup>13</sup>	[13, 64, 93]
DPoS	Medium	Strong	< 51% of validators	<1000	<100 s	Bitshares <sup>14</sup>	[64, 93]
ZKP <sup>15</sup>	Low	Strong	<51% of computing power	—	—	Zcash <sup>16</sup>	[25]

Instances are described along their characteristics, an implementation example, and the source for argument of the claims.

throughput, i.e., how many transactions per second the protocol can successfully process, and (d) the latency, i.e., the time it takes to successfully confirm the transaction.

An interesting observation with respect to consensus protocols in Table 2 is that there seems to be a lack of systematic and empirical studies to test the claims around the proposed consensus protocols and their architectural properties. Of the papers reviewed, none based their statements on these properties and how they behave under different circumstance on the results of empirical testing. An exception is a study by Dihn et al. [25] that provides a framework to benchmark private blockchains. Public blockchains, however, remain untested. Nevertheless, some general observations can be made: With the introduction of PoW, a shift is made from protocols that require high trust among nodes to a more trustless setting. This is not surprising given that one of the main aims of the original Bitcoin protocol was to make it expandable beyond a fixed number of authenticated nodes. Accordingly, an increase in byzantine fault tolerance of consensus protocols can also be observed. However, the results also demonstrate that an increase in the byzantine fault tolerance is coupled with a decrease in throughput and latency. For instance, the Tendermint PoS variant has a higher throughput and lower latency as the Bitcoin PoW protocol, yet it is also less byzantine fault tolerant (<33% to <51%).

<sup>9</sup>[www.hyperledger.org/projects/fabric](http://www.hyperledger.org/projects/fabric).

<sup>10</sup>[www.corda.net](http://www.corda.net).

<sup>11</sup>[www.bitcoin.com](http://www.bitcoin.com).

<sup>12</sup><https://www.hyperledger.org/projects/sawtooth>.

<sup>13</sup>[www.tendermint.com](http://www.tendermint.com).

<sup>14</sup>[www.bitshares.org](http://www.bitshares.org).

<sup>15</sup>ZCash uses an Equihash variant of the PoW protocol; for Equihash the Blocktime differs from that of the Bitcoin PoW no direct data are provided in literature on its throughput and latency.

<sup>16</sup>[www.z.cash](http://www.z.cash).

For the PoET and ZKP proof-based consensus protocols there are no data available with regard to their properties. Moreover, although it is stated in References [18, 25], ZK-SNARK techniques incur large overheads in terms of storage space, and there is little evidence to substantiate these claims.

The gradual shift from PoW toward PoS-based consensus can have some important ramifications for the security of blockchains. With the exception of Reference [53], little efforts have been made to investigate the unresolved security issues of PoS variants and how to resolve address them. One such problem is how to determine the deposit forgers are required to make to participate in the consensus protocol (see Section 5.3.4). If the gains of introducing invalid transactions outweigh the losses of a deposit, then the solution is will not be effective. Sharding of consensus seeking is a novel development mentioned in the papers under review. Yet we view that there are many unaddressed questions concerning blockchain sharding. When sharding a consensus protocol the nodes in the network will be distributed between several shards. How this distribution should take place remains unclear, however. Provided that a secure manner has been found to distribute the nodes, then it should be determined how many nodes must encompass a shard to make it secure as small shards (with few nodes) are easy to attack.

**9.2.3 Blockchain Hybrids Emerging.** Our literature suggests that more recently several interesting combinations of public blockchains with a permissioned model are under experimentation.

We observe from the coding process that the terms permissioned/private blockchain and permissionless/public blockchain are used interchangeably. Based on our extensive review, we propose that blockchain-oriented networks should be categorized based on both (1) network authentication coupled to permissions (permissioned/permissionless) and (2) accessibility of the network (public, consortium, private). Although the first blockchain network (Bitcoin) was public and permissionless, from the coding process it can be observed that in the papers under review permissioned and private/consortium blockchain networks are often mentioned (see online materials in Appendix C). One of the main reasons remarked in literature is that public blockchains have to utilize consensus protocols that can be regarded as slower than the ones that could be employed for private/consortium blockchains. Furthermore, having more control over the permissions each participant in the network has is another reason mentioned. We observe that private/consortium and permissioned networks are often regarded as a substitute to their public counterpart because of these challenges.

**9.2.4 Mainstream Adoption of Blockchain Technology.** An increase of academic works on the topic, and the overall increase of literature produced every year can be observed (see Section 3.5). These results show that BCT has been embraced by the mainstream. The usability of BCT for end-users and developers [60] could, however, hamper the further adoption of BCT. Tools to develop blockchains/smart contracts seem to be absent. Especially for users that aim to employ smart contracts to define simple transaction logic a more simple and user friendly approach that does not involve programming would benefit the mainstream adoption of smart contracts and blockchain.

The results of this research show that governance of blockchain both in terms of the general ecosystem and at the platform level are a challenge currently. In many countries the use of BCT for applications or cryptocurrency has an opaque legal status. Making policy and regulation is difficult, since a blockchain networks operate internationally and are not bound to a single jurisdiction [65, 79]. Implementing updates or changes for a public blockchain network are a platform level governance challenge caused by decentralized decision making involving many participants. A solution to this problem might be the centralization of update permissions. However, that solution would be at odds with decentralized principles of blockchains and introduce several security hazards.

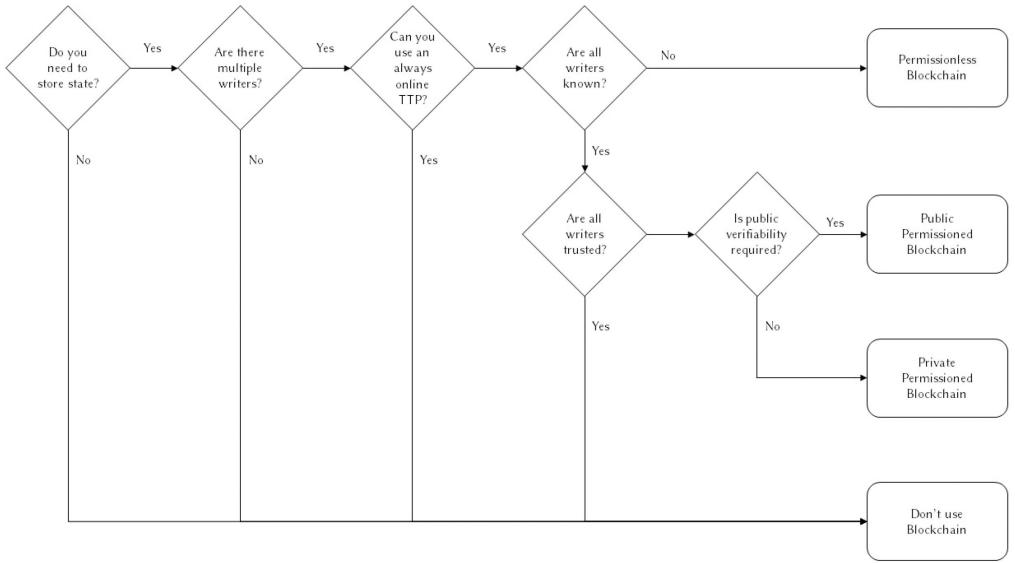


Fig. 11. Decision-making model for blockchain networks; adapted from Reference [87].

**9.2.5 Safe Executability of Legally Binding Smart Contracts.** Smart contracts are an important application for BCT, as can be inferred from the number of times the code has been applied (see Section 9.1). All primary studies focus on discussing the security of Ethereum smart contracts. This trend is highlighted also in related work [2]. The most obvious implication of this shortcoming is that more rigorous, generalisable, and formalized approach to analyzing the safe and secure executability of smart contracts is required. A preliminary investigation of safe concurrent smart-contracts’ executability Dickerson et al. [24] who describe an abstract solution to this problem without any rigorous evaluation. The trend toward sharding blockchain consensus can have implications for the execution of smart contracts what these exactly are remains unknown till date. Conversely, the word “contract” in the smart term definition contract would imply for its legally bound enforceability under specific operational conditions [19].

**9.2.6 Blockchain: Properties versus Challenges.** BCT offers some interesting properties such as disintermediation, programmability, transparency, and verifiability that could potentially be beneficial for many applications (Section 7). However, the technology also has its drawbacks (Section 8). The work of Mougayar [65] provides a decision-making framework based on business parameters to aid in identifying the problems that BCT can streamline. Our work strives to provide a technical perspective to compound the business side explored by works such as Mougayar [65]. Gatteschi et al. [29] suggest that when considering to adopt BCT one should ask (1) whether a shared database is required, (2) whether multiple parties write the data, (3) whether disintermediation is needed, and (4) whether it is required to see the linkage between transactions. In a similar vein, Wüst and Gervais [87] argue that when deciding whether to adopt and design blockchain the contingencies depicted in Figure 11 should be considered. They argue that if a trusted third party (TTP) always is available, then there is no need for a blockchain. Thereafter, one has to make an analysis of the network participants to decide which type of blockchain network is most appropriate.

Besides the considerations presented in References [29, 87] from the results of this study we observe the following. First, a blockchain can disintermediate the transactions made on a P2P network between untrusted participants. However, as a result the throughput of most consensus

protocols for blockchains is low and their latency high (see Table 2). If these properties do not satisfy the requirements of the application, then these consensus protocols are not the most suitable solution. Visa, for instance, handles 2,000 transactions per second [77, 81], while most protocols currently cannot process transactions at that rate. Furthermore, when all network participants can trust each other a consensus protocol is not required. In such situations, other solutions such as distributed databases [69] or P2P network [4] offer a less costly and faster way to exchange data or transactions. Second, not all blockchain platforms allow for the deployment of smart contracts that enable the programmability of transactions. Moreover, currently blockchain-based smart contracts present some unresolved problems such as ensuring security [2] and legal enforceability [19]. In the past three decades, several approaches have been suggested to develop contracts that are executed electronically [45, 51] known as electronic contracts that currently have a higher maturity for these applications. Third, transparency and verifiability are a double-edged sword. In fact, these properties ensure the correctness of the distributed ledger; however, transparency is not always desirable. More transaction privacy could be guaranteed by using ZKP's, yet it is suggested that these incur additional overhead. An alternative is to perform the transactions on a blockchain network with a permissioned model. This could create censorship resistance that jeopardizes the security of the blockchain (see online materials accessible via a link in Appendix F).

### 9.3 Addressing Our Research Questions

9.3.1 *RQ 1: How can blockchain technology be systematically defined.* The first research question for this study concerned how blockchain technology can be defined—the RQ is addressed in five consecutive sections. First, in Section 4 the study defines BCT based on a FCA. The results of the FCA show that the most mentioned attributes of BCT are transactions, chains, distributed ledger, block, P2P network, consensus protocol, hash function, public key, node and private key. The analysis shows that what sets distributed ledger technology apart from BCT is the storage of transactions in blocks. In presenting this definition we present both scholars and practitioners with an understanding of the BCT phenomenon that can provide common ground for discussion. In Section 5.1 describe and flesh out the interplay between the logical elements of BCT. This description aids in understanding how the combination of individual software elements constitute to BCT. The development view presented in Section 5.2 provides DApp developers and platform developers with some guidance what the main components of BCT are. In Section 5.3, we elaborate on the workings of consensus protocols. Essentially, our synthesis of the literature shows that increasing the throughput and decreasing the latency of transactions results in a less secure consensus protocol. By providing this overview practitioners can make more informed decisions about the design of their blockchain platform. Furthermore, the section also pointed out the idiosyncratic problems and lack of identification thereof. Scholars can use this information to focus future research endeavors. We are confident that by creating an understanding of current consensus protocols practitioners will be inspired to improve these protocols or, using the lessons learned, create new consensus protocols with more favorable properties. The arrangement of the P2P network is discussed in Section 5.4, which has important ramifications for the architectural properties of the blockchain. More informed decision-making can aid practitioners in making the optimal choice when adopting blockchain technology.

9.3.2 *RQ 2: What applications of blockchain technology have currently been published and how can these applications be classified?* In Section 6, we present the three main versions of blockchain applications, and we further explain them in online material (See Appendix F). First, we discuss cryptocurrency including interoperability and, briefly, off-chain payment channels. We continue to discuss smart contracts in online material (see Appendix F for access), including current security

issues and trends. In doing so we provide both practitioners and scholars with trends for these two applications. Finally, the online materials (Appendix F) shows a wide array of BCT-based applications. An interesting trend that can be observed is that almost all of these applications use a form of cryptocurrency to exchange assets and smart contracts to automatize processes. The oversight provides an insight into the rich palette of blockchain applications.

**9.3.3 RQ 3: What are the properties of blockchain technology and what are their tradeoffs?** The architecture properties of blockchain technology are discussed in Section 7. The systematic analysis presented in Section 7 couples properties to architectural design. When discussing BCT, the literature generally posits certain properties. However, the analysis shows that these properties depend on the architectural design of a blockchain. Understanding the implications of certain design decisions is therefore crucial. In Section 9.1.1, we discuss the relation between BCT software elements and properties.

**9.3.4 RQ 4: What are the challenges for blockchain technology?** The challenges of BCT are presented in Section 8 along with some directions that are currently being pursued to address these challenges. Section 9.1.5 discusses these challenges and their relative importance to establish a notion of their urgency. The identification of these challenges aid in formulating a future research roadmap and establishes and outlook how they are addressed currently.

In the following section, we address the fourth and final research questions by presenting future opportunities for research.

## 10 RESEARCH GAPS AND ROADMAP

The field of BCT is rapidly developing, yet the results of this study point out four main research gaps that provide opportunities for future research. The research gaps and opportunities for future research have been identified using the results of the GT coding process (Section 9.1) combined with the descriptive (Section 3.5) and in-depth analysis of the papers under review (Section 9.2).

### 10.1 Consensus Protocols

Despite the fact that several publications discuss consensus protocols, few claims regarding the performance of these protocols are substantiated with evidence. One of the research gaps identified by this study is that evidence is lacking, because the consensus protocols suitable for public blockchains remain untested. Similar to the approach presented in Reference [25] for private blockchains, feature research endeavors could *empirically test* the properties of consensus protocols for public blockchains. Addressing this research gap is especially important, since the results of this study suggest that throughput and latency issues are still the major concern, and have a relative high urgency (see Section 9.1.5). Moreover, these empirical results aid practitioners making more informed architectural design choices. There is a trend toward developing consensus protocols based on *sharding* that could be further explored in future research, as with the exception of Reference [55] little work on this topic has been published. In the same vein, a shift toward using PoS-based consensus protocol for public blockchains can be observed (see Section 9.1.2). However, most published works on consensus protocols only specifically discuss PoW or provide a broad overview (see Section 3.5). Only three works [42, 53, 70] have been published on the topic of PoS, and only one of these, Reference [53], addresses the security of PoS protocols. More research on how PoS consensus protocols can effectively ensure the integrity of public blockchains is needed. Especially since there are still open issues to address such as *determining deposits* required for staking [53] during a PoS-based protocol.

## 10.2 Data Storage and Privacy

Privacy has been shown to be a fundamental issue for BCT (see Section 9.1.5). As a remedy zk-SNARKs have been proposed (see Section 5.3.6 and Section 9.1.2). It has been claimed zk-SNARKS incur *large overheads* in terms of storage space. However, there is little evidence to substantiate this claim. Future studies could focus on investigating this claim and, if proven to be true, investigate manners to decrease this overhead. Although zk-SNARKs facilitate private transactions in our sample no studies were found identified that investigate their effects on performance aspects (throughput and latency). Forthcoming research efforts could determine *the impact of zk-SNARKS on performance*, which aids in further developing ZKP protocols in general. There is little work on how blockchains can become made *GDPR compliant* (see Section 8.4). However, we posit that, like any other technology, BCT has to operate within the current legal framework. The increasing amount of *data storage space* required by nodes is an unaddressed challenge. Some efforts have been made to address this problem such as the introduction of lightweight clients. The results of our study show that currently data storage is not perceived as a challenge with a high importance (see Section 9.1.5). Yet we posit that in time the problem might become more urgent when the throughput of consensus protocols will increase and in parallel, the amount of data required to be stored [77]. Thus, future research should investigate more efficient ways to store data or determine if data can be omitted from the ledger.

## 10.3 Smart Contracts

From the results of this study *smart contracts* emerge as an important concept related to BCT (see Section 9.1). All of the works in our sample investigates to smart contracts deployed on the Ethereum blockchain platform. Yet, more rigorous, *generalizable*, and *formalized approaches* to analyzing the safe and secure executability of smart contracts is lacking. Novel research could focus its efforts on ensuring the security and safety of smart contracts for platforms other than Ethereum. Literature related electronic contracts (e.g., References [16, 46, 58]) could be a potential inspiration how to address the security issues of smart contracts. Furthermore, further work is required to establish the viability and implementation of *legally enforceable* smart contracts. Given the trend toward sharding the consensus protocol research should be conducted how *cross-shard smart contract validation* can be performed.

## 10.4 Usability

This study has shown that multiple applications for BCT have been explored. However, the usability of BCT and smart contracts in particular remains limited, since *user-friendly blockchain-oriented tools* are not widely available to non-programmers. The trends in challenges mentioned (see Section 9.1.5) suggest that over time this will become a more pervasive problem. Providing tools and approaches to ease the development of DApps or smart contract could open up BCT to a broader audience. Therefore, *user-friendly tools for smart contract development* would be helpful.

# 11 LIMITATIONS AND THREATS TO VALIDITY

Using the guidelines provided by Wohlin et al. [86] the limitations and threats to validity for this study were identified and are discussed in this section.

**External Validity.** First, developments in the field of BCT are introduced at a fast pace. Hence, some of these developments may exist but have not been published yet. Some papers on the topic are so novel that they have not been indexed yet and not included in the selected items. Another issue is that the terminology is still evolving and universal definitions for concepts such as BCT-based smart contracts have not yet been formalized. This issue has been addressed by including

search terms that are being used interchangeably in the search string (e.g., distributed ledger and blockchain technology) to ensure that all potential aliases of blockchain technology were covered. Items that were found using the search terms have been assessed thoroughly based on various dimensions of quality employing inter-rater reliability. Another threat to validity stems from including aliases of BCT is that findings of the study could also be related to distributed ledger technology. As a strategy to mitigate these threats we set out to define BCT based on a formal concept analysis using literature that specifically mentioned BCT in its introduction and background section.

**Internal and Construct Validity.** Second, to attain the results of the research questions a Glaserian-Straussian GT [30] Grounded-Theory coding approach has been used. Although an inter-rater measurement has been employed, the risk of observer bias is still present. Some additional codes were added to the list established during the pilot study. On these codes, however, no inter-rater assessment has been performed. Furthermore, the current body of knowledge on the topic of BCT to day remains limited. Many proposed consensus protocols in the field of BCT have not been rigorously empirically tested in terms of their properties. A prime example being the PoS consensus protocol for which a thorough assessment of the properties (e.g., throughput and latency) of its many variants is lacking in literature. In the same vein, with the notable exception of Li et al. [53], not much research has been carried out with regard to the security of the PoS variants. In addition, in Section 5.3 we have discussed several consensus protocols found in literature from a process view, indicating their process flow, idiosyncratic security issues and other properties. For the sake of space, only the consensus protocols that were recurring more than 3 times in at least two different papers across our primary studies were included in this study. However, as mentioned, this study has also identified other consensus protocols that were not discussed in full for space sake. Third, the findings of this research are partially based on grey literature sources. Inherent to grey literature is that the quality and accuracy of these sources can be disputable. To mitigate this threat we assessed each grey literature item obtained through our search strategy using multiple criteria based on the guidelines provided by Garousi et al. [28]. Furthermore, the assessment of the grey literature by the first and second authors of this study has been subjected to an inter-rater reliability test (see Section 3.4).

## 12 CONCLUSIONS

This study was enacted to (1) analyze how blockchain technology can be defined, (2) provide a systematic overview of the state-of-the-art concepts around that definition, and (3) distill a grounded research roadmap around the topic. In Section 4, using a Formal Concept Analysis (FCA) [82] approach, a systematic definition of BCT was distilled.

Beyond the operational definition above, using the well-known 4+1 software architecture viewpoint framework [47], the architecture elements of BCT were fleshed out; specifically, our results recap (1) the way a platform can be designed, (2) how transactions are processed, and (3) the architectural arrangements typically used for the P2P network underlying BCT. The third aim of this research was to flesh out what blockchain-based applications have been discussed in the state of the art and how these can be categorized. The study reveals that there are three types of use cases for BCT: (1) cryptocurrencies, (2) smart contracts, and (3) an array of more general applications that can be sub-categorized into five categories, namely (a) security and fraud detection, (b) securities and insurance, (c) record-keeping, (d) Internet-of-Things (IoT), as well as (e) smart legal contracts. As a fourth objective of this study, we set out to determine the architecture properties of blockchain technology and their tradeoffs. Data analysis reveals eight coupled architectural characteristics—these properties are *tradeoffs* exercised during blockchain architectural design. The fifth objective of this research was to identify the challenges for BCT. In the future,

six main challenges for blockchain technology need to be addressed: (1) decreasing latency for the conformation of transactions, (2) increasing the throughput of transactions that is related to the design of the consensus protocol, (3) decreasing data storage requirement, (4) protecting the privacy of blockchain users, (5) data governance of blockchain networks, and (6) the usability of the technology.

Finally, an analysis of the papers under review demonstrates that there are four research gaps that need to be addressed by future research concerning (1) consensus protocols, (2) data privacy and storage, (3) smart contracts, and (4) the usability of blockchain for end users. In our own research agenda, we plan to further analyse the results and data stemming from our study to further provide architectural and decision-making instruments for practitioners and academics alike. Furthermore, we plan to focus around the social and societal concerns around BCT, namely its *privacy-by-design* [32] aspects as well as its end-user acceptance and maturity.

## APPENDIX

### A IN-AND EXCLUSION CRITERIA AND REPLICATION PACKAGE

All in- and exclusion criteria used for literature, data, and resources used in the scope of this study were bundled up and made available for replication purposes. The package in question is available online here: <https://drive.google.com/drive/folders/163pHkUcEHv0tIRKq7MF1pzK0-sM9COGB?usp=sharing>.

### B FORMAL CONCEPT ANALYSIS METHOD

The approach taken for this study to attain the results of our FCA are available at: <https://drive.google.com/drive/folders/1susTjFs8ze9mR957zy8N5--UHQGbkdfN?usp=sharing>.

### C GROUNDED-THEORY ANALYSIS METHOD

For this study, a Grounded-Theory approach has been taken. A complete overview of all the steps in we conducted for our analysis can be found online via <https://drive.google.com/drive/folders/1n4qXkyBns6CW3BXomPgDUtko3uZ0VH3V?usp=sharing>.

### D INTER-RATER RELIABILITY ASSESSMENT METHOD

The results of the coding process have been assessed for inter-rater reliability. The steps, analysis, and the outcomes of the analysis can be found at: [https://drive.google.com/open?id=1WVd\\_7PMvCBLhJ4mTRZrt3cvCRz5wxNy7](https://drive.google.com/open?id=1WVd_7PMvCBLhJ4mTRZrt3cvCRz5wxNy7).

### E A 4+1 OVERVIEW OF BLOCKCHAIN TECHNOLOGY

Further reading based upon the results of this study on the 4+1 view of blockchain technology can be found online: <https://drive.google.com/drive/folders/1lmMq25SCg0LNQ5j7ndONpTtCQJ5Et4O?usp=sharing>.

### F A CATALOGUE OF USECASES FOR BLOCKCHAIN TECHNOLOGY

A more elaborate catalogue of the usecases for blockchain technology is available online via <https://drive.google.com/drive/folders/1fESobEUrkRIreco3moM1lzypaS7jG1LL?usp=sharing>.

### G A GROUNDED THEORY OF BLOCKCHAIN TECHNOLOGY: TRENDS AND FIGURES

All figures and tables related to frequency analysis and trends are available online via [https://drive.google.com/open?id=1FhLn1LS0hV6vWjCx\\_udmKXGsv52b5zM0](https://drive.google.com/open?id=1FhLn1LS0hV6vWjCx_udmKXGsv52b5zM0).

## ACKNOWLEDGMENTS

We thank Omkar Khair for his assistance during this research and his valuable practitioners' perspective on BCT.

## REFERENCES

- [1] 2016. *Embracing Disruption, Tapping the Potential of Distributed Ledger to Improve the Post-Trade Landscape*. Technical Report. New York, NY.
- [2] Maher Alharby and Aad van Moorsel. 2017. A systematic mapping study on current research topics in smart contracts. *Int. J. Comput. Sci. Inf. Technol.* 9, 5 (2017), 151–164.
- [3] Elli Androulaki, Artem Barger, Vita Bortnikov, Christian Cachin, Konstantinos Christidis, Angelo De Caro, David Enyeart, Christopher Ferris, Gennady Laventman, Yacov Manevich, Srinivasan Muralidharan, Chet Murthy, Binh Nguyen, Manish Sethi, Gari Singh, Keith Smith, Alessandro Sorniotti, Chrysoula Stathakopoulou, Marko Vukolić, Sharon Weed Cocco, and Jason Yellick. 2018. Hyperledger fabric: A distributed operating system for permissioned blockchains. In *Proceedings of the 13th European Conference on Computer Systems (EuroSys'18)*. ACM, New York, NY, Article 30, 15 pages. DOI : <https://doi.org/10.1145/3190508.3190538>
- [4] Stephanos Androullos-Theotokis and Diomidis Spinellis. 2004. A survey of peer-to-peer content distribution technologies. *ACM Comput. Surv.* 36, 4 (Dec. 2004), 335–371. DOI : <https://doi.org/10.1145/1041680.1041681>
- [5] Peter Bailis, Arvind Narayanan, Andrew Miller, and Song Han. 2017. Research for practice: Cryptocurrencies, blockchains, and smart contracts; hardware for deep learning. *Commun. ACM* 60, 5 (2017), 48–51.
- [6] Foteini Baldimtsi, Aggelos Kiayias, Thomas Zacharias, and Bingsheng Zhang. 2016. Indistinguishable proofs of work or knowledge. In *Proceedings of the Annual Conference on Advances in Cryptology (ASIACRYPT'16)*, Jung Hee Cheon and Tsuyoshi Takagi (Eds.). Springer, Berlin, 902–933.
- [7] Massimo Bartoletti and Livio Pomianu. 2017. An empirical analysis of smart contracts: Platforms, applications, and design patterns. In *Financial Cryptography and Data Security*, Michael Brenner, Kurt Rohloff, Joseph Bonneau, Andrew Miller, Peter Y. A. Ryan, Vanessa Teague, Andrea Bracciali, Massimiliano Sala, Federico Pintore, and Markus Jakobsson (Eds.). Springer International Publishing, Cham, 494–509.
- [8] L. Bass, P. Clements, and R. Kazman. 2003. *Software Architecture in Practice*. Addison-Wesley. 2003045300
- [9] F. Rizal Batubara, Jolien Ubacht, and Marijn Janssen. 2018. Challenges of blockchain technology adoption for e-government: A systematic literature review. In *Proceedings of the 19th Annual International Conference on Digital Government Research: Governance in the Data Age (dg.o.'18)*. ACM, New York, NY, Article 76, 9 pages. DOI : <https://doi.org/10.1145/3209281.3209317>
- [10] Alex Biryukov, Dmitry Khovratovich, and Ivan Pustogarov. 2014. Deanonymisation of clients in bitcoin P2P network. In *Proceedings of the 2014 ACM SIGSAC Conference on Computer and Communications Security (CCS'14)*. ACM, New York, NY, 15–29. DOI : <https://doi.org/10.1145/2660267.2660379>
- [11] Phillip Boucher. 2017. *How Blockchain Technology Could Change Our Lives*. Technical Report. Brussels, Belgium.
- [12] Richard Gendal Brown, James Carlyle, Ian Grigg, and Mike Hearn. 2016. *Corda: A Distributed Ledger*. Technical Report, New York, NY. Corda.
- [13] Christian Cachin and Marko Vukolić. 2017. Blockchains consensus protocols in the wild. (2017). arXiv:arXiv preprint arXiv:1707.01873
- [14] Davide Calvaresi, Alevtina Dubovitskaya, Jean Paul Calbimonte, Kuldar Taveter, and Michael Schumacher. 2018. Multi-agent systems and blockchain: Results from a systematic literature review. In *Proceedings of the International Conference on Practical Applications of Agents and Multi-Agent Systems*. Springer, Cham, 110–126.
- [15] Fran Casino, Thomas K. Dasaklis, and Constantinos Patsakis. 2019. A systematic literature review of blockchain-based applications: Current status, classification and open issues. *Telematics and Informatics* 36 (2019), 55–81. DOI : <https://doi.org/10.1016/j.telet.2018.11.006>
- [16] Dickson K. W. Chiu, Shing-Chi Cheung, and Sven Till. 2003. A three-layer architecture for e-contract enforcement in an e-service environment. In *Proceedings of the 36th Annual Hawaii International Conference on System Sciences 2003*. IEEE, 10–pp.
- [17] K.-Y. Chow, A. J. David, and A. M. Ionescu-Graff. 1983. Switching capacity relief model: Theoretical development. 1–7. [http://i-teletraffic.org/\\_Resources/Persistent/d1af3fa05b6d58cea28b80af39fd919a290d8b4a/chow831.pdf](http://i-teletraffic.org/_Resources/Persistent/d1af3fa05b6d58cea28b80af39fd919a290d8b4a/chow831.pdf).
- [18] Konstantinos Christidis and Michael Devetsikiotis. 2016. Blockchains and smart contracts for the internet of things. *IEEE Access* 4 (2016), 2292–2303.
- [19] Christopher D. Clack, Vikram A. Bakshi, and Lee Braine. 2016. Smart contract templates: Foundations, design landscape and research directions. (2016). arXiv:1608.00771
- [20] Bram Cohen and Krzysztof Pietrzak. 2018. Simple proofs of sequential work. In *Proceedings of the Annual International Conference on the Theory and Applications of Cryptographic Techniques*. Springer, Cham, 451–467.

- [21] The Ethereum Community. 2017. Ethereum homestead documentation, release 01. (2017). arXiv:<http://www.ethdocs.org/en/latest/>.
- [22] Michael Crosby, Nachi Nachiappan, Pradhan Pattanayak, Sanjeev Verma, and Vignesh Kalyanaraman. 2015. *BlockChain Technology Beyond Bitcoin*. Technical Report. Berkeley, CA.
- [23] Joost de Kruijff and Hans Weigand. 2017. Understanding the blockchain using enterprise ontology. In *Proceedings of the International Conference on Advanced Information Systems Engineering*. Springer, Cham, 29–43.
- [24] Thomas Dickerson, Paul Gazzillo, Maurice Herlihy, and Eric Koskinen. 2017. Adding concurrency to smart contracts. In *Proceedings of the ACM Symposium on Principles of Distributed Computing*. ACM, New York, NY, 303–312.
- [25] Tien Tuan Anh Dinh, Rui Liu, Meihui Zhang, Gang Chen, Beng Chin Ooi, and Ji Wang. 2018. Untangling blockchain: A data processing view of blockchain systems. *IEEE Trans. Knowl. Data Eng.* 30, 7 (2018), 1366–1385.
- [26] Tien Tuan Anh Dinh, Ji Wang, Gang Chen, Rui Liu, Beng Chin Ooi, and Kian-Lee Tan. 2017. Blockbench: A framework for analyzing private blockchains. In *Proceedings of the 2017 ACM International Conference on Management of Data*. ACM, New York, NY, 1085–1100.
- [27] Michael Fröwis and Rainer Böhme. 2017. In code we trust? In *Data Privacy Management, Cryptocurrencies and Blockchain Technology*. Springer, 357–372.
- [28] Vahid Garousi, Michael Felderer, and Mika V. Mäntylä. 2019. Guidelines for including grey literature and conducting multivocal literature reviews in software engineering. *Inf. Softw. Technol.* 106 (2019), 101–121. DOI: <https://doi.org/10.1016/j.infsof.2018.09.006>
- [29] Valentina Gatteschi, Fabrizio Lamberti, Claudio Demartini, Chiara Pranteda, and Víctor Santamaría. 2018. To blockchain or not to blockchain: That is the question. *IT Profess.* 20, 2 (2018), 62–74.
- [30] Barney G. Glaser and Anselm L. Strauss. 2017. *Discovery of Grounded Theory: Strategies for Qualitative Research*. Routledge, New York, NY.
- [31] Florian Glaser. 2017. Pervasive decentralisation of digital infrastructures: A framework for blockchain enabled system and use case analysis. In *Proceedings of the 50th Hawaii International Conference on System Sciences*. 1543–1552.
- [32] Seda Gürses, Carmela Troncoso, and Claudia Diaz. 2011. Engineering privacy by design. *Computers, Privacy & Data Protection* 14, 3 (2011), 25.
- [33] Andrew F. Hayes and Klaus Krippendorff. 2007. Answering the call for a standard reliability measure for coding data. *Commun. Methods Meas.* 1, 1 (2007), 77–89. DOI: <https://doi.org/10.1080/19312450709336664>
- [34] Shuangyu He, Qianhong Wu, Xizhao Luo, Zhi Liang, Dawei Li, Hanwen Feng, Haibin Zheng, and Yanan Li. 2018. A social-network-based cryptocurrency wallet-management scheme. *IEEE Access* 6 (2018), 7654–7663.
- [35] Garick Hileman and Michel Rauchs. 2017. *Global Blockchain Benchmark Study*. Technical Report. Cambridge, United Kingdom.
- [36] W. Hon, Kuan, John Palfreyman, and Matthew Tegart. 2016. *Distributed Ledger Technology and Cybersecurity: Improving Information Security in the Financial Sector*. Technical Report. Heraklion, Greece.
- [37] Aljosha Judmayer, Alexei Zamyatin, Nicholas Stifter, Artemios G. Voyatzis, and Edgar Weippl. 2017. Merged mining: Curse or cure? In *Data Privacy Management, Cryptocurrencies and Blockchain Technology*, Joaquin Garcia-Alfaro, Guillermo Navarro-Arribas, Hannes Hartenstein, and Jordi Herrera-Joancomartí (Eds.). Springer International Publishing, Cham, 316–333.
- [38] Alexey Kalinov. 2006. Scalability of heterogeneous parallel systems. *Program. Comput. Softw.* 32, 1 (2006), 1–7.
- [39] Elena Karafiloski and Anastas Mishev. 2017. Blockchain solutions for big data challenges: A literature review. In *Proceedings of the IEEE 17th International Conference on Smart Technologies (EUROCON'17)*. IEEE, 763–768.
- [40] Rami Khalil and Arthur Gervais. 2017. Revive: Rebalancing off-blockchain payment networks. In *Proceedings of the 2017 ACM SIGSAC Conference on Computer and Communications Security*. ACM, New York, NY, 439–453.
- [41] Minhaj Ahmad Khan and Khaled Salah. 2018. IoT security: Review, blockchain solutions, and open challenges. *Fut. Gener. Comput. Syst.* 82 (2018), 395–411.
- [42] Aggelos Kiayias, Alexander Russell, Bernardo David, and Roman Oliynykov. 2017. Ouroboros: A provably secure proof-of-stake blockchain protocol. In *Proceedings of the Annual Conference on Advances in Cryptology (CRYPTO'17)*, Jonathan Katz and Hovav Shacham (Eds.). Springer International Publishing, Cham, 357–388.
- [43] Barbara Kitchenham. 2004. Procedures for performing systematic reviews. Technical Report. Keele University.
- [44] Kari Korppela, Jukka Hallikas, and Tomi Dahlberg. 2017. Digital supply chain transformation toward blockchain integration. In *Proceedings of the 50th Hawaii International Conference on System Sciences*. IEEE, 4182–4191. DOI: <https://doi.org/10.24251/HICSS.2017.506>
- [45] P. Radha Krishna and Kamalakar Karlapalem. 2012. A methodology for evolving e-contracts using templates. *IEEE Trans. Serv. Comput.* 6, 4 (2012), 497–510.
- [46] P. Radha Krishna, Kamalakar Karlapalem, and Ajay R. Dani. 2005. From contracts to e-contracts: Modeling and enactment. *Inf. Technol. Manage.* 6, 4 (2005), 363–387.
- [47] Philippe B. Kruchten. 1995. The 4+ 1 view model of architecture. *IEEE Softw.* 12, 6 (1995), 42–50.

- [48] Tsung-Ting Kuo, Hugo Zavaleta Rojas, and Lucila Ohno-Machado. 2019. Comparison of blockchain platforms: A systematic review and healthcare examples. *J. Am. Med. Inf. Assoc.* 26, 5 (2019), 462–478.
- [49] Leslie Lamport, Robert Shostak, and Marshall Pease. 1982. The Byzantine generals problem. *ACM Trans. Program. Lang. Syst.* 4, 3 (1982), 382–401.
- [50] Jong-Hyouk Lee. 2018. BiDaaS: Blockchain based ID as a service. *IEEE Access* 6 (2018), 2274–2278.
- [51] Ronald M. Lee. 1988. A logic model for electronic contracting. *Dec. Support Syst.* 4, 1 (1988), 27–44. DOI: [https://doi.org/10.1016/0167-9236\(88\)90096-6](https://doi.org/10.1016/0167-9236(88)90096-6)
- [52] K. Li, H. Li, H. Hou, K. Li, and Y. Chen. 2017. Proof of vote: A high-performance consensus protocol based on vote mechanism & consortium blockchain. In *Proceedings of the 2017 IEEE 19th International Conference on High Performance Computing and Communications; IEEE 15th International Conference on Smart City; IEEE 3rd International Conference on Data Science and Systems (HPCC/SmartCity/DSS'17)*. IEEE, 466–473. DOI: <https://doi.org/10.1109/HPCC-SmartCity-DSS.2017.61>
- [53] Wenting Li, Sébastien Andreina, Jens-Matthias Bohli, and Ghassan Karamé. 2017. Securing proof-of-stake blockchain protocols. In *Data Privacy Management, Cryptocurrencies and Blockchain Technology*. Springer, Cham, 297–315.
- [54] Loi Luu, Duc-Hiep Chu, Hrishi Olickel, Prateek Saxena, and Aquinas Hobor. 2016. Making smart contracts smarter. In *Proceedings of the 2016 ACM SIGSAC Conference on Computer and Communications Security (CCS'16)*. ACM, New York, NY, 254–269. DOI: <https://doi.org/10.1145/2976749.2978309>
- [55] Loi Luu, Viswesh Narayanan, Chaodong Zheng, Kunal Baweja, Seth Gilbert, and Prateek Saxena. 2016. A secure sharding protocol for open blockchains. In *Proceedings of the 2016 ACM SIGSAC Conference on Computer and Communications Security*. ACM, New York, NY, 17–30.
- [56] Loi Luu, Jason Teutsch, Raghav Kulkarni, and Prateek Saxena. 2015. Demystifying incentives in the consensus computer. In *Proceedings of the 22nd ACM SIGSAC Conference on Computer and Communications Security (CCS'15)*. ACM, New York, NY, 706–719. DOI: <https://doi.org/10.1145/2810103.2813659>
- [57] Daniel Macrinici, Cristian Cartofeanu, and Shang Gao. 2018. Smart contract applications within blockchain technology: A systematic mapping study. *Telemat. Inf.* 35, 8 (2018), 2337–2354. DOI: <https://doi.org/10.1016/j.tele.2018.10.004>
- [58] Olivera Marjanovic and Zoran Milosevic. 2001. Towards formal modeling of e-contracts. In *Proceedings of the 5th IEEE International Enterprise Distributed Object Computing Conference*. IEEE, 59–68.
- [59] Ivan Martinovic, Lucas Kello, and Ivo Slagavonic. 2017. *Blockchains for Governmental Services: Design Principles, Applications, and Case Studies*. Technical Report. Centre for Technology and Global Affairs, University of Oxford, Oxford, United Kingdom.
- [60] Jan Mendling, Ingo Weber, Wil Van Der Aalst, Jan Vom Brocke, Cristina Cabanillas, Florian Daniel, Søren Debois, Claudio Di Ciccio, Marlon Dumas, Schahram Dustdar, Avigdor Gal, Luciano García-Bañuelos, Guido Governatori, Richard Hull, Marcello La Rosa, Henrik Leopold, Frank Leymann, Jan Recker, Manfred Reichert, Hajo A. Reijers, Stefanie Rinderle-Ma, Andreas Solti, Michael Rosemann, Stefan Schulte, Munindar P. Singh, Tjits Slaats, Mark Staples, Barbara Weber, Matthias Weidlich, Mathias Weske, Xiwei Xu, and Liming Zhu. 2018. Blockchains for business process management - Challenges and opportunities. *ACM Trans. Manage. Inf. Syst.* 9, 1, Article 4 (Feb. 2018), 16 pages. DOI: <https://doi.org/10.1145/3183367>
- [61] Weizhi Meng, Elmar Wolfgang Tischhauser, Qingju Wang, Yu Wang, and Jinguang Han. 2018. When intrusion detection meets blockchain technology: A review. *IEEE Access* 6 (2018), 10179–10188.
- [62] Andrew Miller, Iddo Bentov, Ranjit Kumaresan, and Patrick McCorry. 2017. Sprites: Payment channels that go faster than lightning. *arXiv preprint arXiv:1702.05812* (2017).
- [63] David Mills, Kathy Wang, Brendan Malone, Anjana Ravi, Jeff Marquardt, Clinton Chen, Anton Badev, Timothy Brezinski, Linda Fahy, Kimberley Liao, Vanessa Kargenian, Max Ellithorpe, Wendy Ng, and Maria Baird. 2016. *Distributed Ledger Technology in Payments, Clearing, and Settlement*. Technical Report. Federal Reserve United States, Washington, DC.
- [64] Du Mingxiao, Ma Xiaofeng, Zhang Zhe, Wang Xiangwei, and Chen Qijun. 2017. A review on consensus algorithm of blockchain. In *Proceedings of the IEEE International Conference on Systems, Management and Cybernetics (SMC'17)*. IEEE, 2567–2572.
- [65] William Mougayar. 2016. *The Business Blockchain: Promise, Practice, and Application of the Next Internet Technology*. John Wiley & Sons.
- [66] Satoshi Nakamoto. 2008. Bitcoin: A peer-to-peer electronic cash system. (2008).
- [67] Harish Natarajan, Karla Krause, Solvej, and Luskin Gradstein, Helen. 2017. *Distributed Ledger Technology (DLT) and Blockchain*. Technical Report. World Bank, Washington, DC.
- [68] Svein Ølnes, Jolien Ubacht, and Marijn Janssen. 2017. Blockchain in government: Benefits and implications of distributed ledger technology for information sharing. *Gov. Inf. Quart.* 34, 11 (2017), 355–364.
- [69] M. Tamer Özsu and Patrick Valduriez (Eds.). 2011. *Principles of Distributed Database Systems* (3 ed.). Springer, New York.

- [70] Rafael Pass and Elaine Shi. 2017. The sleepy model of consensus. In *Proceedings of the International Conference on the Theory and Application of Cryptology and Information Security*. Springer, Cham, 380–409.
- [71] M. Pease, R. Shostak, and L. Lamport. 1980. Reaching agreement in the presence of faults. *J. ACM* 27, 2 (Apr. 1980), 228–234. DOI : <https://doi.org/10.1145/322186.322188>
- [72] Thierry Perroud and Reto Inversini. 2013. *Enterprise Architecture Patterns: Practical Solutions for Recurring IT-Architecture Problems*. Springer, Berlin. DOI : <https://doi.org/10.1007/978-3-642-37561-3>
- [73] Gareth W. Peters and Efstatios Panayi. 2015. Understanding modern banking ledgers through blockchain technologies: Future of transaction processing and smart contracts on the internet of money. (2015). arXiv:1511.05740
- [74] Marc Pilkington. 2016. *Research Handbook on Digital Transformations*. Edward Elgar Publishing, Cheltenham, UK.
- [75] Ric Shreves. 2017. *A Revolution in Trust: Distributed Ledger Technology in Relief and Development*. Technical Report. Mercy Corps, Portland, OR.
- [76] Josh Stark. 2017. *Applications of Distributed Ledger Technology to Regulatory and Compliance Processes*. Technical Report. R3, New York, NY.
- [77] Melanie Swan. 2015. *Blockchain: Blueprint for a New Economy*. O'Reilly Media, Inc., Sebastopol, CA.
- [78] Bayu Adhi Tama, Bruno Joachim Kweka, Youngho Park, and Kyung-Hyune Rhee. 2017. A critical review of blockchain and its current applications. In *Proceedings of the International Conference on Electrical Engineering and Computer Science (ICECOS'17)*. IEEE, 109–113.
- [79] Don Tapscott and Alex Tapscott. 2017. *Realizing the Potential of Blockchain, A Multi Stakeholder Approach to the Stewardship of Blockchain and Cryptocurrencies*. Technical Report. Geneva, Switzerland.
- [80] Florian Tschorsch and Björn Scheuermann. 2016. Bitcoin and beyond: A technical survey on decentralized digital currencies. *IEEE Commun. Surv. Tutor.* 18, 3 (2016), 2084–2123.
- [81] Jeroen van Oerle and Patrick Lemmens. 2016. *Distributed Ledger technology for the Financial Industry*. Technical Report. Amsterdam, The Netherlands.
- [82] Frano Škopjanac Maćina and Bruno Blašković. 2014. Formal concept analysis, overview and applications. *Proc. Eng.* 69 (2014), 1258–1267.
- [83] Mark Walport. 2016. *Distributed Ledger Technology: Beyond Blockchain*. Technical Report. London, United Kingdom.
- [84] Yingli Wang, Jeong Hugh Han, and Paul Beynon-Davies. 2019. Understanding blockchain technology for future supply chains: A systematic literature review and research agenda. *Supply Chain Manage.* 24, 1 (2019), 62–84.
- [85] Ingo Weber, Xiwei Xu, Régis Riveret, Guido Governatori, Alexander Ponomarev, and Jan" Mendling. 2016. Untrusted business process monitoring and execution using blockchain. In *Business Process Management*, Marcello La Rosa, Peter Loos, and Oscar Pastor (Eds.). Springer International Publishing, Cham, 329–347.
- [86] Claes Wohlin, Per Runeson, Martin Höst, Magnus C Ohlsson, Björn Regnell, and Anders Wesslén. 2012. *Experimentation in Software Engineering*. Springer Science & Business Media, Cham.
- [87] Wüst and Arthur Gervais. 2017. Do you need a blockchain? (2017). DOI : <https://doi.org/20180828:104906> arXiv:<https://eprint.iacr.org/2017/375>
- [88] Xiwei Xu, Cesare Pautasso, Liming Zhu, Vincent Gramoli, Alexander Ponomarev, An Binh Tran, and Shiping Chen. 2016. The blockchain as a software connector. In *Proceedings of the 2016 13th Working IEEE/IFIP Conference on Software Architecture (WICSA'16)*. IEEE, 182–191.
- [89] X. Xu, I. Weber, M. Staples, L. Zhu, J. Bosch, L. Bass, C. Pautasso, and P. Rimba. 2017. A taxonomy of blockchain-based systems for architecture design. In *Proceedings of the 2017 IEEE International Conference on Software Architecture (ICSA'17)*. IEEE, 243–252. DOI : <https://doi.org/10.1109/ICSA.2017.33>
- [90] Dylan Yaga, Peter Mell, Nik Roby, and Karen Scarfone. 2018. *Blockchain Technology Overview*. Technical Report. Gaithersburg, MA.
- [91] Jesse Yli-Huumo, Deokyoon Ko, Sujin Choi, Sooyong Park, and Kari Smolander. 2016. Where is current research on blockchain technology? A systematic review. *PLoS ONE* 11, 10 (10 2016), 1–27. DOI : <https://doi.org/10.1371/journal.pone.0163477>
- [92] Fei Richard Yu, Jianmin Liu, Ying He, Pengbo Si, and Yanhua Zhang. 2018. Virtualization for distributed ledger technology (vDLT). *IEEE Access* 6 (2018), 25019–25028.
- [93] Z. Zheng, S. Xie, H. Dai, X. Chen, and H. Wang. 2017. An overview of blockchain technology: Architecture, consensus, and future trends. In *Proceedings of the IEEE International Congress on Big Data (BigData Congress'17)*. IEEE, 557–564. DOI : <https://doi.org/10.1109/BigDataCongress.2017.85>
- [94] Aviv Zohar. 2015. Bitcoin: Under the hood. *Commun. ACM* 58, 9 (2015), 104–113.

Received April 2019; revised September 2019; accepted October 2019